



---

# Lecture 7

## Basics for Machine Learning and A Special Emphasis on CNN

Lin ZHANG, PhD  
School of Software Engineering  
Tongji University  
Spring 2022



# 10 Breakthrough Technologies 2017 (MIT Tech Review)



Reversing Paralysis



Self Driving




Paying with  
Your Face



Practical Quantum  
Computers



The 360-  
Degree Selfie




Hot Solar Cells



Gene Therapy 2.0



The Cell Atlas



Botnets of Things



Reinforcement  
Learning

The core is machine learning



傍晚，小街路面上沁出微雨后的湿润，和煦的细风吹来，抬头看看天边的晚霞，嗯，明天又是一个好天气。走到水果摊旁，挑了个根蒂蜷缩、敲起来声音浊响的青绿西瓜，一边满心期待着皮薄肉厚瓢甜的爽落感，一边愉快地想着：这学期狠下了功夫，基础概念弄得清清楚楚，算法作业也是信手拈来，这门课成绩一定差不了！

摘自《机器学习》（周志华著，2016）



# Outline

---

- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection




# A little history about AI

## 人工智能

1956年，**麦卡锡**召集哈佛大学、麻省理工学院、IBM公司、贝尔实验室的研究人员召开达特茅斯会议正式提出“人工智能”



2006年达特茅斯会议当事人重聚，左起：**摩尔**、**麦卡锡**、**明斯基**、**赛弗里奇**、**所罗门诺夫**



John McCarthy

**人工智能之父**

**人工智能**是指计算机系统具备的能力，该能力可以履行原本只有依靠人类智慧才能完成的复杂任务





# A little history about AI

---

## 人工智能

### 🔍 人工智能的最终目标

探讨智能形成的基本机理，研究利用自动机模拟人的思维过程



### 📍 人工智能的近期目标

研究如何使计算机去做那些靠人的智力才能做的工作



# A little history about AI

## 人工智能的研究范式及历程

**符号主义**: 采用知识表达和逻辑符号系统来模拟人类的智能，试图对智能进行宏观研究 (**Knowledge-driven**)

**联接主义**: 始于W.S. McCulloch和皮兹 (Pitts) 的先驱工作，直到目前的深度学习，是微观意义上的探索 (**Data driven**)

**生物启发的智能**: 依赖于生物学、脑科学、生命科学和心理学等学科的发现，将机理变为可计算的模型 (**Biology mechanism driven**)

1950-1960

1960-1970

1970-1980

1990-2015

2016后

二者独立并驾齐驱

符号主义：专家系统和知识工程为主流


符号主义滞步，日本第五代计算机失败，联接主义蓬勃发展

联接主义占据主导；同时模糊逻辑取得重大进展

生物启发的智能—跨模态的信息处理



# A little history about AI



# A little history about AI

## 人工智能产业发展加速明显


### 自然语言处理 (NLP) :

微软 Skype Translator 实现同声传译



### 计算机视觉 (CV) :

Face++ 的人脸识别云服务



### 计算机视觉 (CV) :

格林深瞳的视频监控可智能识别犯罪



### 感知、规划和决策:

Google 无人驾驶汽车





# A little history about AI

---

## 人工智能成为世界焦点



人工智能目前已经成为世界各国关注的焦点。2017年7月，中国政府发布了“新一代人工智能发展规划”

✓ 人工智能是开启未来智能世界的秘钥，是未来科技发展的战略制高点；谁掌握人工智能，谁就将成为未来核心技术的掌控者



# What is machine learning?

---

- Gives "computers the ability to learn without being explicitly programmed" (Arthur Samuel in 1959)



Arthur Lee Samuel  
(December 5, 1901 – July 29, 1990)

- It explores the study and construction of algorithms that can learn from and make predictions on data
- It is employed in a range of computing tasks where designing and programming explicit algorithms with good performance is difficult or unfeasible

[1] Samuel, Arthur L., Some Studies in Machine Learning Using the Game of Checkers, IBM Journal of Research and Development, 1959



# Supervised VS Unsupervised

---

- Supervised learning
  - It will infer a function from labeled training data
  - The training data consists of a set of training examples
  - Each example is a pair consisting of an input object (typically a vector) and a desired output value (also called the supervisory signal)
- Unsupervised learning
  - Trying to find hidden structure in unlabeled data
  - Since the examples given to the learner are unlabeled, there is no error or reward signal to evaluate a potential solution
  - Such as PCA, K-means (a clustering algorithm)



# About sample

---

- Attribute (feature), attribute value, label, and example



{青绿, 蜷缩, 浊响: 好瓜}  
feature values      label value  
one example



# Training, testing, and validation

---

- Training sample and training set

A training set comprising  $m$  training samples,

$$D = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m)\}$$

where  $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{id}) \in \chi$  is the feature vector of  $i$ th sample and  $y_i \in \psi$  is its label

By training, our aim is to find a mapping,

$$f : \chi \mapsto \psi$$

based on  $D$

If  $\psi$  comprises discrete values, such a prediction task is called “**classification**”; if it comprises real numbers, such a prediction task is called “**regression**”



# Training, testing, and validation

---

- Training sample and training set
- Test set
  - A test set is a set of data that is independent of the training data, but that follows the same probability distribution as the training data
  - Used only to assess the performance of a fully specified classifier



# Training, testing, and validation

---

- Training sample and training set
- Test set
- Validation set
  - In order to avoid overfitting, when any classification parameter needs to be adjusted, it is necessary to have a validation set; it is used for model selection
  - The training set is used to train the candidate algorithms, while the validation set is used to compare their performances and decide which one to take



# Overfitting, Generalization, and Capacity

---

- Overfitting
  - It occurs when a statistical model describes random error or noise instead of the underlying relationship
  - It generally occurs when a model is excessively complex, such as having too many parameters relative to the number of observations
  - A model that has been overfit will generally have poor predictive performance, as it can exaggerate minor fluctuations in the data



# Overfitting, Generalization, and Capacity

---


- Overfitting
- Generalization
  - Refers to the performance of the learned model on new, previously unseen examples, such as the test set




# Overfitting, Generalization, and Capacity

- Overfitting
- Generalization


Example: Linear regression (housing prices)



$\rightarrow \theta_0 + \theta_1 x$   
"Underfit" "High bias"



$\rightarrow \theta_0 + \theta_1 x + \theta_2 x^2$   
"Just right"




$\rightarrow \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$   
<http://Overfit.csdn.net/Highvariance09>

# Overfitting, Generalization, and Capacity

- Overfitting
- Generalization


Example: Logistic regression




$$h_{\theta}(x) = g(\underline{\theta_0 + \theta_1 x_1 + \theta_2 x_2})$$

( $g$  = sigmoid function)

*"Underfit"*



$$g(\underline{\theta_0 + \theta_1 x_1 + \theta_2 x_2} + \underline{\theta_3 x_1^2} + \underline{\theta_4 x_2^2} + \underline{\theta_5 x_1 x_2})$$



$$g(\underline{\theta_0 + \theta_1 x_1 + \theta_2 x_1^2} + \underline{\theta_3 x_1^2 x_2} + \underline{\theta_4 x_1^2 x_2^2} + \underline{\theta_5 x_1^2 x_2^3} + \underline{\theta_6 x_1^3 x_2} + \dots)$$

<http://blog.osdn.net/zouxy09>



# Overfitting, Generalization, and Capacity

---


- Overfitting
- Generalization
- Capacity
  - Measures the complexity, expressive power, richness, or flexibility of a classification algorithm
  - Ex, DCNN (deep convolutional neural networks) is powerful since its capacity is very large

$$y^* = b + \omega x, \quad y^* = b + \omega_1 x_1 + \omega_2 x_2, \quad y^* = b + \sum_{i=1}^{10} \omega_i x_i$$

→ higher capacity



# Overfitting, Generalization, and Capacity



→ higher capacity



# Performance Evaluation

---

Given a sample set (training, validation, or test)

$$D = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m)\}$$

To assess the performance of the learner  $f$ , we need to compare the prediction  $f(\mathbf{x})$  and its ground-truth label  $y$

For regression task, the most common performance measure is MSE (mean squared error),

$$E(f; D) = \frac{1}{m} \sum_{i=1}^m (f(\mathbf{x}_i) - y_i)^2$$



# Performance Evaluation (for classification)

---

- Error rate
  - The ratio of the number of misclassified samples to the total number of samples

$$E(f; D) = \frac{1}{m} \sum_{i=1}^m \mathbf{1}(f(\mathbf{x}_i) \neq y_i)$$

- Accuracy
  - It is derived from the error rate

$$acc(f; D) = \frac{1}{m} \sum_{i=1}^m \mathbf{1}(f(\mathbf{x}_i) = y_i) = 1 - E(f; D)$$



# Performance Evaluation (for classification)

- Precision and Recall

Ground truth	Prediction	
	positive	negative
positive	True Positive (TP)	False Negative (FN)
negative	False Positive (FP)	True Negative (TN)

$$\textit{precision} = \frac{TP}{TP + FP}$$

$$\textit{recall} = \frac{TP}{TP + FN}$$




# Performance Evaluation (for classification)

---

- Precision and Recall

- Often, there is an inverse relationship between precision and recall, where it is possible to increase one at the cost of reducing the other
- Usually, PR-curve is not monotonic





# Performance Evaluation (for classification)

---

- Precision-recall should be used together; it is meaningless to use only one of them
- However, in many cases, people want to know explicitly which algorithm is better; we can use  $F$ -measure

$$F_{\beta} = \frac{(1 + \beta^2) \times P \times R}{(\beta^2 \times P) + R}$$



# Performance Evaluation (for classification)

- To derive a single performance measure

Varying threshold, we can have a series of  $(P, R)$  pairs,

$$(P_1, R_1), (P_2, R_2), \dots, (P_n, R_n)$$

Then,

$$P_{macro} = \frac{1}{n} \sum_{i=1}^n P_i \quad R_{macro} = \frac{1}{n} \sum_{i=1}^n R_i$$

$$F_{\beta-macro} = \frac{(1 + \beta^2) \times P_{macro} \times R_{macro}}{(\beta^2 \times P_{macro}) + R_{macro}}$$



# Class-imbalance Issue

---

- Problem definition
  - It is the problem in machine learning where the total number of a class of data is far less than the total number of another class of data
  - This problem is extremely common in practice
- Why is it a problem?
  - Most machine learning algorithms work best when the number of instances of each classes are roughly equal
  - When the number of instances of one class far exceeds the other, problems arise



# Class-imbalance Issue

---

- How to deal with this issue?
  - Modify the cost function
  - Under-sampling, throwing out samples from majority classes
  - Oversampling, creating new virtual samples for minority classes
    - » Just duplicating the minority classes could lead the classifier to overfitting to a few examples
    - » Instead, use some algorithm for oversampling, such as SMOTE (synthetic minority over-sampling technique)<sup>[1]</sup>

[1] N.V. Chawla *et al.*, SMOTE: Synthetic Minority Over-sampling Technique, *J. Artificial Intelligence Research* 16: 321-357, 2002



# Class-imbalance Issue

- Minority oversampling by SMOTE<sup>[1]</sup>

Add new minority class instances by:

- For each minority class instance c
  - neighbours = Get KNN(5)
  - n = Random pick one from neighbours
  - Create a new minority class r instance using c's feature vector and the feature vector's difference of n and c multiplied by a random number
    - » i.e.  $r.\text{feats} = c.\text{feats} + (n.\text{feats} - c.\text{feats}) * \text{rand}(0,1)$

[1] N.V. Chawla *et al.*, SMOTE: Synthetic Minority Over-sampling Technique, J. Artificial Intelligence Research 16: 321-357, 2002



# Outline

---

- Basic concepts
- Linear model
  - Linear regression
  - Logistic regression
  - Softmax regression
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection



# Linear regression

---

- Our goal in linear regression is to predict a target continuous value  $y$  from a vector of input values  $\mathbf{x} \in R^d$ ; we use a linear function  $h$  as the model
- At the training stage, we aim to find  $h(\mathbf{x})$  so that we have  $h(\mathbf{x}_i) \approx y_i$  for each training sample  $(\mathbf{x}_i, y_i)$
- We suppose that  $h$  is a linear function, so

$$h_{(\theta,b)}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x} + b, \boldsymbol{\theta} \in R^{d \times 1}$$

Rewrite it,

$$\boldsymbol{\theta}' = \begin{pmatrix} \boldsymbol{\theta} \\ b \end{pmatrix}, \mathbf{x}' = \begin{pmatrix} \mathbf{x} \\ 1 \end{pmatrix}$$

$$\boldsymbol{\theta}'^T \mathbf{x}' + b = \boldsymbol{\theta}'^T \mathbf{x}' \equiv h_{\boldsymbol{\theta}'}(\mathbf{x}')$$

Later, we simply use  $h_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}, \boldsymbol{\theta} \in R^{(d+1) \times 1}, \mathbf{x} \in R^{(d+1) \times 1}$



# Linear regression

---

- Then, our task is to find a choice of  $\theta$  so that  $h_{\theta}(x_i)$  is as close as possible to  $y_i$

The cost function can be written as,

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (\theta^T x_i - y_i)^2$$

Then, the task at the training stage is to find

$$\theta^* = \arg \min_{\theta} \frac{1}{2} \sum_{i=1}^m (\theta^T x_i - y_i)^2$$

For this special case, it has a closed-form optimal solution

Here we use a more general method, **gradient descent** method



# Linear regression

---

- Gradient descent
  - It is a first-order optimization algorithm
  - To find a local minimum of a function, one takes steps proportional to the negative of the gradient of the function at the current point
  - One starts with a guess  $\theta_0$  for a local minimum of  $J(\theta)$  and considers the sequence such that

$$\theta_{n+1} := \theta_n - \alpha \nabla_{\theta} J(\theta)_{|\theta=\theta_n}$$

where  $\alpha$  is called as **learning rate**



# Linear regression

---

- Gradient descent






# Linear regression

---

- Gradient descent





# Linear regression

- Gradient descent

Repeat until convergence ( $J(\theta)$  will not reduce anymore)

{

$$\theta_{n+1} := \theta_n - \alpha \nabla_{\theta} J(\theta)_{|\theta=\theta_n}$$

}

GD is a general optimization solution; for a specific problem,  
the key step is how to compute gradient



# Linear regression

---

- Gradient of the cost function of linear regression

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (\theta^T \mathbf{x}_i - y_i)^2$$

The gradient is,

$$\nabla_{\theta} J(\theta) = \begin{bmatrix} \frac{\partial J(\theta)}{\partial \theta_1} \\ \frac{\partial J(\theta)}{\partial \theta_2} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_{d+1}} \end{bmatrix} \text{ where, } \frac{\partial J(\theta)}{\partial \theta_j} = \sum_{i=1}^m (h_{\theta}(\mathbf{x}_i) - y_i) x_{ij}$$



# Linear regression

---

- Some variants of gradient descent
  - The ordinary gradient descent algorithm looks at every sample in the **entire** training set on every step; it is also called as **batch gradient descent**
  - **Stochastic gradient descent (SGD)** repeatedly run through the training set, and each time when we encounter a training sample, we update the parameters according to the gradient of the error w.r.t that single training sample only

Repeat until convergence

{

    for  $i = 1$  to  $m$  ( $m$  is the number of training samples)

    {

$$\theta_{n+1} := \theta_n - \alpha (\theta_n^T x_i - y_i) x_i$$

    }

}



# Linear regression

---

- Some variants of gradient descent
  - The ordinary gradient descent algorithm looks at every sample in the **entire** training set on every step; it is also called as **batch gradient descent**
  - **Stochastic gradient descent (SGD)** repeatedly run through the training set, and each time when we encounter a training sample, we update the parameters according to the gradient of the error w.r.t that single training sample only
  - **Minibatch SGD**: it works identically to SGD, except that it uses more than one training samples to make each estimate of the gradient



# Linear regression

---

- More concepts
  - $m$  Training samples can be divided into  $N$  minibatches
  - When the training sweeps all the batches, we say we complete one **epoch** of training process; for a typical training process, several epochs are usually required

```
epochs = 10;  
numMiniBatches = N;  
while epochIndex < epochs && not convergent  
{  
    for minibatchIndex = 1 to numMiniBatches  
    {  
        update the model parameters based on this minibatch  
    }  
}
```



# Outline

---

- Basic concepts
- Linear model
  - Linear regression
  - Logistic regression
  - Softmax regression
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection



# Logistic regression

---

- Logistic regression is used for binary classification
- It squeezes the linear regression  $\theta^T x$  into the range (0, 1) ; thus the prediction result can be interpreted as probability
- At the testing stage


The probability that the testing sample  $x$  is positive is represented as  $h_\theta(x) = \frac{1}{1 + \exp(-\theta^T x)}$

The probability that the testing sample  $x$  is negative is represented as  $1 - h_\theta(x)$

Function  $\sigma(z) = \frac{1}{1 + \exp(-z)}$  is called as **sigmoid** or **logistic** function



# Logistic regression



The shape of sigmoid function

One property of the sigmoid function

$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$



*Can you verify?*



# Logistic regression

---

- The hypothesis model can be written neatly as

$$P(y | \mathbf{x}; \boldsymbol{\theta}) = (h_{\boldsymbol{\theta}}(\mathbf{x}))^y (1 - h_{\boldsymbol{\theta}}(\mathbf{x}))^{1-y}$$

- Our goal is to search for a value  $\boldsymbol{\theta}$  so that  $h_{\boldsymbol{\theta}}(\mathbf{x})$  is large when  $\mathbf{x}$  belongs to “1” class and small when  $\mathbf{x}$  belongs to “0” class

Thus, given a training set with binary labels  $\{(\mathbf{x}_i, y_i) : i = 1, \dots, m\}$ , we want to maximize,

$$\prod_{i=1}^m (h_{\boldsymbol{\theta}}(\mathbf{x}_i))^{y_i} (1 - h_{\boldsymbol{\theta}}(\mathbf{x}_i))^{1-y_i}$$

Equivalent to maximize,

$$\sum_{i=1}^m y_i \log(h_{\boldsymbol{\theta}}(\mathbf{x}_i)) + (1 - y_i) \log(1 - h_{\boldsymbol{\theta}}(\mathbf{x}_i))$$



# Logistic regression

---

- Thus, the cost function for the logistic regression is (we want to minimize),

$$J(\theta) = -\sum_{i=1}^m y_i \log(h_\theta(x_i)) + (1 - y_i) \log(1 - h_\theta(x_i))$$

To solve it with gradient descent, gradient needs to be computed,

$$\nabla_\theta J(\theta) = \sum_{i=1}^m x_i (h_\theta(x_i) - y_i)$$






# Logistic regression

---

- Exercise
  - Use logistic regression to perform digital classification





# Outline

---

- Basic concepts
- Linear model
  - Linear regression
  - Logistic regression
  - Softmax regression
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection



# Softmax regression

---

- Softmax operation
  - It squashes a  $K$ -dimensional vector  $\mathbf{z}$  of arbitrary real values to a  $K$ -dimensional vector  $\sigma(\mathbf{z})$  of real values in the range  $(0, 1)$ . The function is given by,

$$\sigma(\mathbf{z})_j = \frac{\exp(z_j)}{\sum_{k=1}^K \exp(z_k)}$$

- Since the components of the vector  $\sigma(\mathbf{z})$  sum to one and are all strictly between 0 and 1, they represent a categorical probability distribution



# Softmax regression

---

- For multiclass classification, given a test input  $x$ , we want our hypothesis to estimate  $p(y = k \mid x)$  for each value  $k=1,2,\dots,K$



# Softmax regression

---

- The hypothesis should output a  $K$ -dimensional vector giving us  $K$  estimated probabilities. It takes the form,

$$h_{\phi}(\mathbf{x}) = \begin{bmatrix} p(y=1 | \mathbf{x}; \phi) \\ p(y=2 | \mathbf{x}; \phi) \\ \vdots \\ p(y=K | \mathbf{x}; \phi) \end{bmatrix} = \frac{1}{\sum_{j=1}^K \exp\left(\left(\boldsymbol{\theta}_j\right)^T \mathbf{x}\right)} \begin{bmatrix} \exp\left(\left(\boldsymbol{\theta}_1\right)^T \mathbf{x}\right) \\ \exp\left(\left(\boldsymbol{\theta}_2\right)^T \mathbf{x}\right) \\ \vdots \\ \exp\left(\left(\boldsymbol{\theta}_K\right)^T \mathbf{x}\right) \end{bmatrix}$$

where  $\phi = [\boldsymbol{\theta}_1, \boldsymbol{\theta}_2, \dots, \boldsymbol{\theta}_K] \in R^{(d+1) \times K}$



# Softmax regression

---

- In softmax regression, for each training sample we have,

$$p(y_i = k \mid \mathbf{x}_i; \boldsymbol{\phi}) = \frac{\exp\left(\left(\boldsymbol{\theta}_k\right)^T \mathbf{x}_i\right)}{\sum_{j=1}^K \exp\left(\left(\boldsymbol{\theta}_j\right)^T \mathbf{x}_i\right)}$$

At the training stage, we want to maximize  $p(y_i = k \mid \mathbf{x}_i; \boldsymbol{\phi})$  for each training sample for the correct label  $k$



# Softmax regression

---

- Cost function for softmax regression

$$J(\phi) = -\sum_{i=1}^m \sum_{k=1}^K 1\{y_i = k\} \log \frac{\exp\left(\left(\theta_k\right)^T \mathbf{x}_i\right)}{\sum_{j=1}^K \exp\left(\left(\theta_j\right)^T \mathbf{x}_i\right)}$$

where  $1\{.\}$  is an indicator function

- Gradient of the cost function

$$\nabla_{\theta_k} J(\phi) = -\sum_{i=1}^m \left[ \mathbf{x}_i \left( 1\{y_i = k\} - p(y_i = k | \mathbf{x}_i; \phi) \right) \right]$$



*Can you verify?*



# Cross entropy

---

- After the softmax operation, the output vector can be regarded as a discrete probability density function
- For multiclass classification, the ground-truth label for a training sample is usually represented in one-hot form, which can also be regarded as a density function  
For example, we have 10 classes, and the  $i$ th training sample belongs to class 7, then  $y_i = [0 \ 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0]$
- Thus, at the training stage, we want to minimize

$$\sum_i dist(h(\mathbf{x}_i; \boldsymbol{\theta}), y_i)$$

How to define  $dist$ ? Cross entropy is a common choice



# Cross entropy

- Information entropy is defined as the average amount of information produced by a probabilistic stochastic source of data  $H(X) = -\sum_i p(x_i) \log p(x_i)$



Claude Shannon  
1916–2001

After graduating from Michigan and MIT, Shannon joined the AT&T Bell Telephone laboratories in 1941. His paper ‘A Mathematical Theory of Communication’ published in the *Bell System Technical Journal* in 1948 laid the foundations for modern information theory.

This paper introduced the word ‘bit’, and his concept that information could be sent as a stream of 1s and 0s paved the way for the communications revolution. It is said that von Neumann recommended to Shannon that he use the term entropy, not only because of its similarity to the quantity used in physics, but also because “nobody knows what entropy really is, so in any discussion you will always have an advantage”.



# Cross entropy

---

- Information entropy is defined as the average amount of information produced by a probabilistic stochastic source of data  $H(X) = -\sum_i p(x_i) \log p(x_i)$
- Cross entropy can measure the difference between two distributions

$$H(p, q) = -\sum_i p(x_i) \log q(x_i)$$

where  $p$  is the ground truth and  $q$  is the prediction result

- For multiclass classification, the last layer usually is a softmax layer and the loss is the ‘cross entropy’



# Cross entropy

---

Example:

Suppose that the label of one sample is [1 0 0]

For model 1, the output of the last softmax layer is [0.5 0.4 0.1]

Its cross entropy is (base 10),

$$H(p, q) = -\sum_i p(x_i) \log q(x_i) = -(1 * \log 0.5 + 0 * \log 0.4 + 0 * \log 0.1) \approx 0.3$$

For model 2, the output of the last softmax layer is [0.8 0.1 0.1]

$$H(p, q) = -\sum_i p(x_i) \log q(x_i) = -(1 * \log 0.8 + 0 * \log 0.1 + 0 * \log 0.1) \approx 0.1$$

Model 2 is better



# Outline

---


- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection




# Neural networks

---

- Neural networks are proposed by Walter Harry Pitts and Warren Sturgis McCulloch in 1943



Walter H. Pitts, 23 April  
1923 – 14 May 1969, USA



Warren Sturgis McCulloch, Nov. 16,  
1898 – Sep. 24, 1969, USA



# Neural networks

- In 1948, McCulloch concerned that eventually AI might rule humankind, a topic of much concern in these increasingly automated times. The below article from the September 22, 1948 Brooklyn Daily Eagle records his clarion call about the future

## Warns Machines May Some Day Take Over World

Pasadena, Cal., Sept. 22 (U.P.)—A bearded psychiatrist predicts that man may build machines with complex mechanical minds and space ships may explore the planets independently of human direction.

What's more, said Dr. Warren S. McCulloch, professor of psychiatry at the University of Illinois Medical School, "when the machines become really complex, they may also become neurotic if badgered by frustration in solving problems."

At the same time man learns more about the mind, he is also learning to build thought machines so efficient they may some day take over civilization. McCulloch warned a symposium on cerebral mechanisms at California Institute of Technology.

Some "electronic brains" already have temporary memory banks, and as more efficient vacuum tubes are devised the machines will acquire permanent "memories" to endow them with judgment, he claimed.

Special circuits will give them curiosity and an "instinct" for self-preservation, he predicted.



# Neural networks

---


- It is one way to solve a supervised learning problem given labeled training examples  $\{\mathbf{x}_i, y_i\} (i = 1, \dots, m)$
- Neural networks give a way of defining a complex, non-linear form of hypothesis  $h_{W,b}(\mathbf{x})$ , where  $W$  and  $b$  are the parameters we need to learn from training samples

# Neural networks

---

- A single neuron

- $x_1, x_2$ , and  $x_3$  are the inputs,  $+1$  is the intercept term,  $h_{W,b}(x)$  is the output of this neuron



$$h_{W,b}(x) = f(W^T x) = f\left(\sum_{i=1}^3 W_i x_i + b\right)$$


where  $f(\cdot)$  is the activation function

# Neural networks

---

- Commonly used activation functions
  - Sigmoid function

$$f(z) = \frac{1}{1 + \exp(-z)}$$






# Neural networks

---

- Commonly used activation functions
  - Tanh function

$$f(z) = \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$






# Neural networks

---

- Commonly used activation functions
  - Rectified linear unit (ReLU)

$$f(z) = \max(0, z)$$

$$f(u) = \max(0, u)$$






# Neural networks

---

- Commonly used activation functions
  - Leaky Rectified linear unit (ReLU)

$$f(z) = \begin{cases} z, & \text{if } z > 0 \\ 0.01z, & \text{otherwise} \end{cases}$$






# Neural networks


- Commonly used activation functions
  - Softplus (can be regarded as a smooth approximation to ReLU)

$$f(z) = \ln(1 + e^z)$$



# Neural networks

- A neural network is composed by hooking together many simple neurons
- The output of a neuron can be the input of another
- Example, a three layers neural network,





# Neural networks

---

- Terminologies about the neural network
  - The circle labeled +1 are called **bias units**
  - The leftmost layer is called the **input layer**
  - The rightmost layer is the **output layer**
  - The middle layer of nodes is called the **hidden layer**
    - » In our example, there are 3 input units, 3 hidden units, and 1 output unit
  - We denote the activation (output value) of unit  $i$  in lay  $l$  as
$$a_i^{(l)}$$



# Neural networks

---

$$a_1^{(2)} = f\left(W_{11}^{(1)}x_1 + W_{12}^{(1)}x_2 + W_{13}^{(1)}x_3 + b_1^{(1)}\right)$$


$$a_2^{(2)} = f\left(W_{21}^{(1)}x_1 + W_{22}^{(1)}x_2 + W_{23}^{(1)}x_3 + b_2^{(1)}\right)$$

$$a_3^{(2)} = f\left(W_{31}^{(1)}x_1 + W_{32}^{(1)}x_2 + W_{33}^{(1)}x_3 + b_3^{(1)}\right)$$

$$h_{W,b}(\mathbf{x}) = a_1^{(3)} = f\left(W_{11}^{(2)}a_1^{(2)} + W_{12}^{(2)}a_2^{(2)} + W_{13}^{(2)}a_3^{(2)} + b_1^{(2)}\right)$$

# Neural networks

- Neural networks can have multiple outputs
- Usually, we can add a softmax layer as the output layer to perform multiclass classification





# Neural networks

---

- At the testing stage, given a test input  $x$ , it is straightforward to evaluate its output
- At the training stage, given a set of training samples, we need to train  $W$  and  $b$ 
  - The key problem is how to compute the gradient
  - Backpropagation algorithm



# Neural networks

---


- Backpropagation
  - A common method of training artificial neural networks and used in conjunction with an optimization method such as gradient descent
  - Its purpose is to compute the partial derivative of the loss to each parameter (weights)
  - neural nets will be very large: impractical to write down gradient formula by hand for all parameters
  - recursive application of the chain rule along a computational graph to compute the gradients of all parameters



# Neural networks

---


- Backpropagation



# Neural networks


---

- Backpropagation




# Neural networks

- Backpropagation




# Neural networks

- Backpropagation




# Neural networks

- Backpropagation



# Neural networks

- Backpropagation





# Outline

---

- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- CNN for object detection



# A little history about deep learning

---

## 深度学习



2006年，Hinton和他的学生Salakhutdinov在《科学》上发表了一篇文章，开启了深度学习在学术界和工业界的浪潮。

Reducing the Dimensionality of Data with Neural Networks

文章提出了深层网络训练中梯度消失问题的解决方案：  
**无监督预训练对权值进行初始化+有监督训练微调**

2012年，Hinton课题组为了证明深度学习的潜力，首次参加ImageNet 图像识别比赛，其通过构建的CNN网络AlexNet一举夺得冠军，且碾压第二名（SVM方法）的分类性能。也正是由于该比赛，CNN吸引到了众多研究者的注意



# A little history about deep learning



Yann LeCun



Geoffrey Hinton




Yoshua Bengio

2018 ACM A.M. Turing Award for conceptual and engineering breakthroughs that have made deep neural networks a critical component of computing



# A little history about deep learning

## 深度学习的发展







# A little history about deep learning

## 深度学习的应用领域


语音识别




人脸识别



图像分类




自动驾驶



Deep Learning



生命科学



游戏博弈



语言翻译



# A little history about deep learning

## 深度学习为什么有效



### 开源的计算平台

Caffe、Tensorflow、Pytorch 等





### 计算力是引擎

GPU服务器、集群



### 数据就是燃料

如ImageNet图像分类数据集等  
百万张图片






# A little history about deep learning

## 深度学习助力人工智能

深度学习的“层次化”学习思想与人类视觉认知机理高度适应

1958年，约翰霍普金斯大学 David Hubel 和 Torsten Wiesel 发现人的视觉系统的信息处理是分级的，人对物品的识别可能是一个不断迭代不断抽象的过程。




这一发现是神经科学与认知领域的重大突破，促进了人工智能领域以后五十年的发展



# A little history about deep learning


## 深度学习助力人工智能





# A little history about deep learning

## 深度学习助力人工智能




与传统人工智能相比，深度学习的最大特点就在于能够自动学习与任务相适应的特征



# A little history about deep learning

---

## 深度学习与人工智能



人工智能是一个非常广泛的问题，机器学习是解决这类问题的一个重要手段。深度学习则是机器学习的一个分支。在很多人工智能问题上，深度学习的方法突破了传统机器学习方法的瓶颈，推动了人工智能领域的发展



# When you know these terminologies, you are an expert of DL

Conv layer  
Pooling  
padding  
Stride  
Relu  
Fully connection  
Batch normalization  
Skip

LeNet  
AlexNet  
NIN  
GoogLeNet  
VGGNet  
ResNet  
DenseNet  
MobileNet  
ShuffleNet  
PeleeNet  
ShuffleNetV2  
EfficientNet

RCNN  
Fast-RCNN  
Faster-RCNN  
SPPNet  
MaskNet  
SSD  
Yolo  
YoloV2  
YoloV3  
YoloV4  
Pelee-SSD

Caffe  
Caffe2  
TensorFlow  
MatConvNet  
Theano  
Torch  
PyTorch  
Keras  
MXNet  
DIGITS  
TensorRT



# When you know these terminologies, you are an expert of DL

---

GTX980  
GTX980Ti  
GTX1080  
GTX1080Ti  
TitanX  
TitanXP  
Titan RTX  
Tesla K20  
Tesla K40  
Tesla K80  
Jetson TX1  
Jetson TX2  
PX2

ImageNet  
VOC2007  
VOC2012  
MS COCO  
Kitti

Gradient vanishing  
Adversarial Samples

Geoffrey Hinton  
Alex Krizhevsky  
Yoshua Bengio  
Yann LeCun  
Ian Goodfellow  
Kaiming He  
Ross Girshick  
Joseph Redmon

SYNTHIA  
Virtual Kitti  
URSA  
VIPER  
synMT  
Grand Theft Auto V



# Convolutional neural network

---

- Specially designed for data with grid-like structures  
(LeCun et al. 98)
  - 1D grid: sequential data
  - 2D grid: image
  - 3D grid: video, 3D image volume
- Beat all the existing computer vision technologies on object recognition on ImageNet challenge with a large margin in 2012



# Convolutional neural network

---

- Something you need to know about DCNN
  - Traditional model for PR: fixed/engineered features + trainable classifier
  - For DCNN: it is usually an **end-to-end** architecture; learning data representation and classifier together
  - The learned features from big datasets are **transferable**
  - For training a DCNN, usually we use a **fine-tuning** scheme
  - For training a DCNN, to avoid overfitting, **data augmentation** can be performed



# Convolutional neural network


---

- Problems of fully connected networks
  - Every output unit interacts with every input unit
  - The number of weights grows largely with the size of the input image
  - Pixels in distance are less correlated



# Convolutional neural network

- Problems of fully connected networks



Ranzato CVPR'13



# Convolutional neural network


---

- One simple solution is locally connected neural networks
  - Sparse connectivity: a hidden unit is only connected to a local patch (weights connected to the patch are called filter or kernel)
  - It is inspired by biological systems, where a cell is sensitive to a small sub-region of the input space, called a receptive field; Many cells are tiled to cover the entire visual field



# Convolutional neural network

- One simple solution is locally connected neural networks




Example: 1000x1000 image  
1M hidden units  
Filter size: 10x10  
100M parameters

Ranzato CVPR'13

# Convolutional neural network


---

- One simple solution is locally connected neural networks
  - The learned filter is a spatially local pattern
  - A hidden node at a higher layer has a larger receptive field in the input
  - Stacking many such layers leads to “filters” (not anymore linear) which become increasingly “global”



# Convolutional neural network

- The first CNN
  - LeNet<sup>[1]</sup>



**CNN called LeNet by Yann LeCun (1998)**

[1] Y. LeCun et al., Gradient-based Learning Applied to Document Recognition, Proceedings of the IEEE, Vol. 86, pp. 2278-2324, 1998



# Convolutional neural network

---




Yann LeCun (1960-)  
New York University  
Facebook Artificial Intelligence Research  
A founding father of convolutional nets

(His name is a little difficult to pronounce)



# Convolutional neural network




He gives himself a Chinese name, 杨立昆 (Mar. 2017)



# Convolutional neural network

- Convolution
  - Computing the responses at hidden nodes is equivalent to convoluting the input image  $x$  with a learned filter  $w$

$$net[i, j] = (x * w)[i, j] = \sum_m \sum_n x[m, n]w[i - m, j - n]$$





# Convolutional neural network

---


- Downsampled convolution layer (optional)
  - To reduce computational cost, we may want to skip some positions of the filter and sample only every  $s$  pixels in each direction. A downsampled convolution function is defined as

$$net(i, j) = (\mathbf{x} * \mathbf{w})[i \times s, j \times s]$$

- $s$  is referred as the stride of this downsampled convolution
  - Also called as strided convolution

# Convolutional neural network

- Multiple filters
  - Multiple filters generate multiple feature maps
  - Detect the spatial distributions of multiple visual patterns





Ranzato CVPR'13

Lin ZHANG, SSE, 2022



# Convolutional neural network

- 3D filtering when input has multiple feature maps




Ranzato CVPR'13



# Convolutional neural network

- Convolutional layer



Ranzato CVPR'13



# Convolutional neural network


---

- To the convolution responses, we then perform nonlinear activation
  - ReLU
  - Tanh
  - Sigmoid
  - Leaky ReLU
  - Softplus

# Convolutional neural network

- Local contrast normalization (optional)
  - Normalization can be done within a neighborhood along both spatial and feature dimensions

$$h_{i+1,x,y,k} = \frac{h_{i,x,y,k} - m_{i,N(x,y,k)}}{\sigma_{i,N(x,y,k)}}$$





# Convolutional neural network


---

- Then, we perform pooling
  - Max-pooling partitions the input image into a set of rectangles, and for each sub-region, outputs the maximum value
  - Non-linear down-sampling
  - The number of output maps is the same as the number of input maps, but the resolution is reduced
  - Reduce the computational complexity for upper layers and provide a form of translation invariance
  - Average pooling can also be used



# Convolutional neural network

- Then, we perform pooling




Ranzato CVPR'13



# Convolutional neural network



- Typical architecture of CNN
  - Convolutional layer increases the number of feature maps
  - Pooling layer decreases spatial resolution
  - LCN and pooling are optional at each stage



Ranzato CVPR'13

# Convolutional neural network

- Typical architecture of CNN





Example with only two filters.

Ranzato CVPR'13

# Convolutional neural network

- Typical architecture of CNN

One stage (zoom)





A hidden unit in the first hidden layer is influenced by a small neighborhood (equal to size of filter).

Ranzato CVPR'13

# Convolutional neural network

- Typical architecture of CNN




After a few stages, residual spatial resolution is very small.

We have learned a descriptor for the whole image. **Ranzato CVPR'13**

# Convolutional neural network

- Typical architecture of CNN



Convolution

Pooling



# Convolutional neural network

---

- Some notes about the CNN layers in most recent net architectures
  - Spatial pooling (such as max pooling) is not recommended now. It is usually replaced by a strided convolution, allowing the network to learn its own spatial downsampling
  - Fully connected layers are not recommended now; instead, the last layer is replaced by global average pooling (for classification problems, the number of feature map channels of the last layer should be the same as the number of classes)



# Convolutional neural network

---

- Example:
  - Train a digit classification model (LeNet) and then test it



Finish this exercise in lab session



# Convolutional neural network


---

- Opensource platforms for CNN
  - CAFFE official, <http://caffe.berkeleyvision.org/>
  - Tensorflow, <https://www.tensorflow.org/>
  - Pytorch, [www.pytorch.org/](http://www.pytorch.org/)
  - Theano, <http://deeplearning.net/software/theano/>



# Convolutional neural network

- An online tool for network architecture visualization
  - <http://ethereon.github.io/netscope/quickstart.html>
  - Network architecture conforms to the CAFFE prototxt format
  - The parameter settings and the output dimension of each layer can be conveniently observed





# Outline

---

- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
  - AlexNet
  - NIN
  - GoogLeNet
  - ResNet
  - DenseNet
  - PeleeNet
- CNN for object detection



# AlexNet (NIPS 2012)

---

- AlexNet: CNN for object recognition on ImageNet challenge
  - Trained on one million images of 1000 categories collected from the web with two GPU. 2GB RAM on each GPU. 5GB of system memory
  - Training lasts for one week
  - Google and Baidu announced their new visual search engines with the same technology six months after that
  - Google observed that the accuracy of their visual search engine was doubled

[1] A. Krizhevsky *et al.*, ImageNet classification with deep convolutional neural networks, in Proc. NIPS, 2012



# AlexNet (NIPS 2012)


- ImageNet
  - <http://www.image-net.org/>





# AlexNet (NIPS 2012)


- Architecture of AlexNet
  - 5 convolutional layers and 2 fully connected layers for learning features
  - Max-pooling layers follow first, second, and fifth convolutional layers





# AlexNet (NIPS 2012)

- Architecture of AlexNet
  - The first time deep model is shown to be effective on large scale computer vision task
  - The first time a very large scale deep model is adopted
  - GPU is shown to be very effective on this large deep model





# Network In Network (NIN, ICLR 2014)


---

- Main idea of NIN
  - Conventional convolutional layers uses linear filters followed by a nonlinear activation function to abstract the information within a receptive field
  - Instead, NIN uses micro neural networks with more complex structures to abstract the data within the receptive field
  - The feature maps are obtained by sliding the micro network over the input in a similar manner as CNN
  - Moreover, they use global average pooling over feature maps in the classification layer, which is easier to interpret and less prone to overfitting than traditional fully connected layers


[1] M. Liu *et al.*, Network in network, in Proc. ICLR, 2014



# Network In Network (NIN, ICLR 2014)



(a) Linear convolution layer




(b) Mlpconv layer  
(the last layer of MLP has  $n$  nodes)

- Comparison of linear convolution layer and mlpconv layer
  - Both the layers map the local receptive field to an output feature vector
  - The mlpconv layer maps the input local patch to the output feature vector with a multilayer perceptron (MLP) consisting of multiple fully connected layers with nonlinear activation functions



# Network In Network (NIN, ICLR 2014)



The overall structure of NIN. The last layer is the global average pooling

- More about global average pooling
  - Fully connected layers are prone to overfitting
  - If there are  $c$  classes, the last MLP layer should output  $c$  feature maps, one feature map for each corresponding category of the classification task
  - Take the average of each feature map to get a  $c$  dimensional vector for softmax classification



# Network In Network (NIN, ICLR 2014)

- NIN can be implemented with conventional convolutional layers


For a mlpconv layer, suppose that the input feature map is of the size  $m \times m \times 32$ , the expected output feature map is of the size  $m \times m \times 64$ , the receptive field is  $5 \times 5$ ; the mlpconv layer has 2 hidden layers, whose node numbers are 16 and 32, respectively.

*How to implement this mlpconv layer with convolutional layers?*



# GoogLeNet (CVPR 2015)

- Main idea: make the network deeper and wider, while keeping the number of parameters
- Inception module



[1] C. Szegedy *et al.*, Going deeper with convolutions, in Proc. CVPR, 2015



# GoogLeNet (CVPR 2015)

---

- Many Inception modules can stack together to form a very deep network
- GoogLeNet refers to the version the authors submitted for the ILSVRC 2014 competition
  - This network consists 27 layers (including pooling layers)



# ResNet (CVPR 2016 Best Paper)

- What is the problem of stacking more layers using conventional CNNs?
  - Vanishing gradient, which can hamper the convergence
  - Accuracy get saturated, and then degraded





Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error.

[1] K. He *et al.*, Deep residual learning for image recognition, in Proc. CVPR, 2016




# ResNet (CVPR 2016 Best Paper)

Is there any better way to design deeper networks?  
Answer: Residual learning



Conventional CNN



Residual block



# ResNet (CVPR 2016 Best Paper)

---

- It is easier to optimize the residual mapping ( $F(\mathbf{x})$ ) than to optimize the original mapping ( $H(\mathbf{x})$ )
- Identity mapping is implemented by shortcut
- A residual learning block is defined as,

$$\mathbf{y} = F(\mathbf{x}, \{W_i\}) + \mathbf{x}$$

where  $\mathbf{x}$  and  $\mathbf{y}$  are the input and output vectors of the layers  
 $F + \mathbf{x}$  is performed by a shortcut connection and element-wise addition

Note: If the dimensions of  $F$  and  $\mathbf{x}$  are not equal (usually caused by changing the numbers of input and output channels), a linear projection  $W_s$  (implemented with  $1*1$  convolution) is performed on  $\mathbf{x}$  to match the dimensions,

$$\mathbf{y} = F(\mathbf{x}, \{W_i\}) + W_s \mathbf{x}$$



# ResNet (CVPR 2016 Best Paper)

---

- It is easier to optimize the residual mapping ( $F(\mathbf{x})$ ) than to optimize the original mapping ( $H(\mathbf{x})$ )
- Identity mapping is implemented by shortcut
- A residual learning block is defined as,

$$\mathbf{y} = F(\mathbf{x}, \{W_i\}) + \mathbf{x}$$

where  $\mathbf{x}$  and  $\mathbf{y}$  are the input and output vectors of the layers  
 $F + \mathbf{x}$  is performed by a shortcut connection and element-wise addition

I highly recommend you to take a look the prototxt file of ResNet (<https://github.com/KaimingHe/deep-residual-networks>)



# DenseNet (CVPR 2017 Best Paper)


---

- Highly motivated by ResNet
- A DenseNet comprises “dense blocks” and transition layers
  - Within a dense block, connect all layers with each other in a feed-forward fashion
  - In contrast to ResNet, DenseNet combine features by concatenating them
  - The number of output feature maps of each layer is set as a constant within a dense block and is called as “growth rate”
  - Between two blocks, there is a transition layer, consisting of batch normalization,  $1 \times 1$  convolution, and average pooling

[1] G. Huang *et al.*, Densely connected convolutional networks, in Proc. CVPR, 2017

# DenseNet (CVPR 2017 Best Paper)

- Highly motivated by ResNet
- A DenseNet comprises “dense blocks” and transition layers




A sample dense block, whose growth rate  
is 4



# DenseNet (CVPR 2017 Best Paper)

- Highly motivated by ResNet
- A DenseNet comprises “dense blocks” and transition layers



A sample DenseNet with three dense blocks



# DenseNet (CVPR 2017 Best Paper)

---

- Highly motivated by ResNet
- A DenseNet comprises “dense blocks” and transition layers
- More details about DenseNet design
  - Bottleneck layers. A  $1 \times 1$  convolution layer can be introduced as bottleneck layer before each  $3 \times 3$  convolution to reduce the number of input feature maps, and thus to improve computational efficiency
  - Compression. If a dense block contains  $m$  feature maps, we let the following transition layer generate  $\theta m$  output feature maps where  $0 < \theta \leq 1$  is referred to as the compression factor



# PeleeNet (ICLR Workshop 2018)

---

- MobileNet and ShuffleNet are designed to run on mobile devices with limited computing power and memory resource
- However, they both depend on depthwise separable convolution which lacks efficient implementation in most deep learning frameworks
- PeleeNet
  - based on conventional convolution instead
  - follows the innovative connectivity pattern and some of key design principles of DenseNet
  - achieves a higher accuracy by 0.6% and 11% lower computational cost than MobileNet
  - PeleeNet is only 66% of the model size of MobileNet


[1] J. Wang *et al.*, Pelee: A Real-Time Object Detection System on Mobile Devices, in Proc. ICLR Workshop, 2018

# PeleeNet (ICLR Workshop 2018)

- PeleeNet's design

- Two-Way Dense Layer

Motivated by GoogLeNet, use a 2-way dense layer to get different scales of receptive fields. One way of the layer uses a small kernel size ( $3 \times 3$ ), which is good enough to capture small-size objects. The other way of the layer uses two stacked  $3 \times 3$  convolution to learn visual patterns for large objects






# PeleeNet (ICLR Workshop 2018)

- PeleeNet's design

- Two-Way Dense Layer
- Stem block

The stem block can effectively improve the feature expression ability without adding computational cost too much - better than other more expensive methods, e.g., increasing channels of the first convolution layer



stem block



# PeleeNet (ICLR Workshop 2018)

---

- PeleeNet's design

- Two-Way dense layer
- Stem block
- Dynamic number of channels in bottleneck layer
- Transition layer without compression
- Composite Function

To improve actual speed, they use the conventional wisdom of post activation (Convolution - Batch Normalization - Relu) as their composite function instead of pre-activation used in DenseNet



# Outline

---

- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- CNN for object detection



# Background


---

- Detection is different from classification
  - An image classification problem is predicting the label of an image among the predefined labels; It assumes that there is **single** object of interest in the image and it covers a significant portion of image
  - Detection is about not only finding the class of object but also localizing the extent of an object in the image; the object can be lying anywhere in the image and can be of any size (scale)



# Background

---



Multiple objects detection



# Background

---

- Traditional methods of detection involved using a block-wise orientation histogram (SIFT or HOG) feature which could not achieve high accuracy in standard datasets such as PASCAL VOC; these methods encode a very low level characteristics of the objects and therefore are not able to distinguish well among the different labels
- Deep learning based methods have become the state-of-the-art in object detection in image; they construct a representation in a hierarchical manner with increasing order of abstraction from lower to higher levels of neural network



# Background

---

- Recent developments of CNN based object detectors
  - R-CNN (CVPR 2014)
  - Fast-RCNN (ICCV 2015)
  - Faster-RCNN (NIPS 2015)
  - Yolo (CVPR 2016)
  - SSD (ECCV 2016)
  - Yolov2 (CVPR 2017)
  - YoloV3
  - Pelee-SSD (ICLR Workshop 2018)



- Brute-force idea
  - One could perform detection by carrying out a classification on different sub-windows or patches or regions extracted from the image. The patch with high probability will not only the class of that region but also implicitly gives its location too in the image
  - One brute force method is to run classification on all the sub-windows formed by sliding different sized patches (to cover each and every location and scale) all through the image

Quite Slow!!!




# R-CNN

---

- R-CNN therefore uses an object proposal algorithm (selective search) in its pipeline which gives out a number (~2000) of TENTATIVE object locations
- These object regions are warped to fixed sized (227X227) regions and are fed to a classification convolutional network which gives the individual probability of the region belonging to background and classes



## Region-based Convolution Networks (R-CNNs)



Input  
image


Extract region  
proposals (~2k / image)  
e.g., selective search  
[van de Sande, Uijlings et al.]

Compute CNN  
features on  
regions

Classify and refine  
regions

# Fast-RCNN

- Compared to RCNN
  - A major change is a single network with two loss branches pertaining to soft-max classification and bounding box regression
  - This multitask objective is a salient feature of Fast-RCNN as it no longer requires training of the network independently for classification and localization





# Faster-RCNN

---

- Compared to Fast-RCNN
  - Faster-RCNN replaces Selective Search with CNN itself for generating the region proposals (called RPN-region proposal network) which gives out tentative regions at almost negligible amount of time




# Yolo (YoloV4) is a state-of-the-art method

---

- Yolo (You Only Look Once)
  - The major exceptional idea is that it tackles the object detection problem as a regression problem
  - A single neural network predicts bounding boxes and class probabilities directly from full images in one evaluation
  - The whole detection pipeline is a single network
  - It is extremely fast; With TitanX, it can process ~50 frames/s



# Yolo (YoloV4) is a state-of-the-art method





# Yolo (YoloV4) is a state-of-the-art method

GitHub - AlexeyAB/darknet: YOLO v4 | [2004.10934] YOLOv4: Optimal S... | CVPR 2021 Open Access Repository | +  
→ ⌂ ⌂ 🔒 https://github.com/AlexeyAB/darknet  
video\_yolov4.sh Added sample scripts for yolov4 2 years ago

☰ README.md

## Yolo v4, v3 and v2 for Windows and Linux

### (neural networks for object detection)

Paper YOLO v4: <https://arxiv.org/abs/2004.10934>

Paper Scaled YOLO v4: \* [CVPR 2021](#): use to reproduce results: [ScaledYOLOv4](#)

More details in articles on medium:

- [Scaled\\_YOLOv4](#)
- [YOLOv4](#)

Manual: <https://github.com/AlexeyAB/darknet/wiki>

Discussion:

- [Reddit](#)
- [Google-groups](#)
- [Discord](#)

