

Machine Learning

KNN

Dr. Shuang LIANG

Recall: Logistic Regression

- Model

$$f_{w,b}(x) = \sigma \left(\sum_i w_i x_i + b \right)$$

Output: between 0 and 1

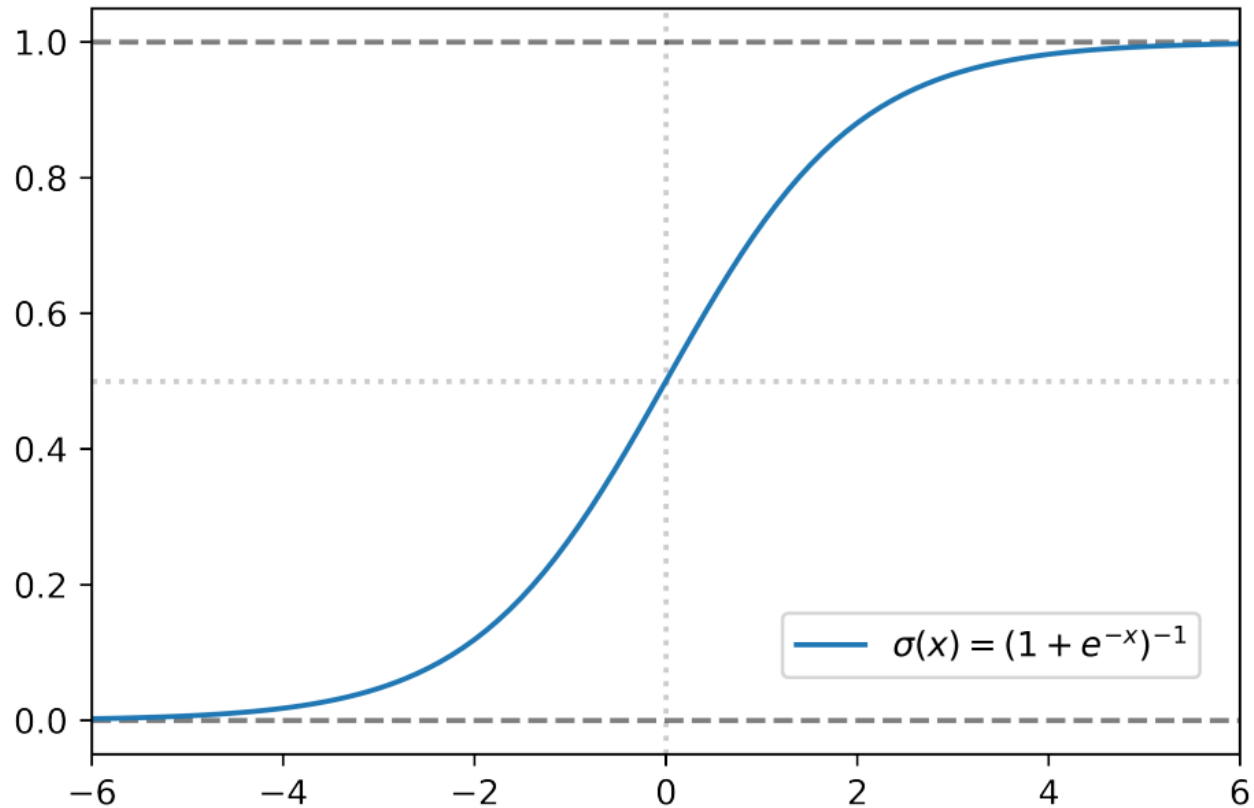
- Loss: Cross Entropy

$$= \sum_n - \left[\hat{y}^n \ln f_{w,b}(x^n) + (1 - \hat{y}^n) \ln (1 - f_{w,b}(x^n)) \right]$$

- Optimization: Gradient Descent

$$w_i \leftarrow w_i - \eta \sum_n \underbrace{\left(\hat{y}^n - f_{w,b}(x^n) \right)}_{\text{error}} x_i^n$$

Recall: Sigmoid



Today's Topics

- Type of classifiers
- KNN
- Setting Parameters
- Analysis of KNN

Today's Topics

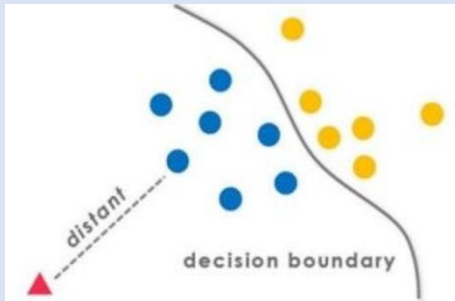
- *Type of classifiers*
- KNN
- Setting Parameters
- Analysis of KNN

Types of Classifiers

Model-based

Discriminative

directly estimate a decision rule/boundary



Logistic regression

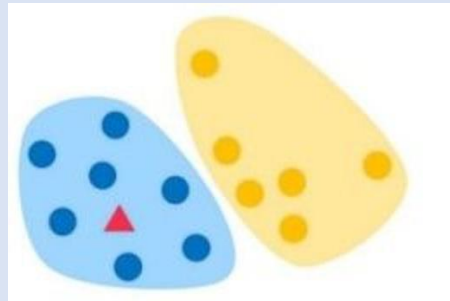
Decision tree

Neural network

.....

Generative

build a generative statistical model



Naïve Bayes

Bayesian Networks

HMM

.....

No Model

Instance-based

Use observation directly

KNN

Discriminative

- Only care about estimating the conditional probabilities $P(y|x)$
- Very good when underlying distribution of data is really complicated (e.g. texts, images, movies)

Generative

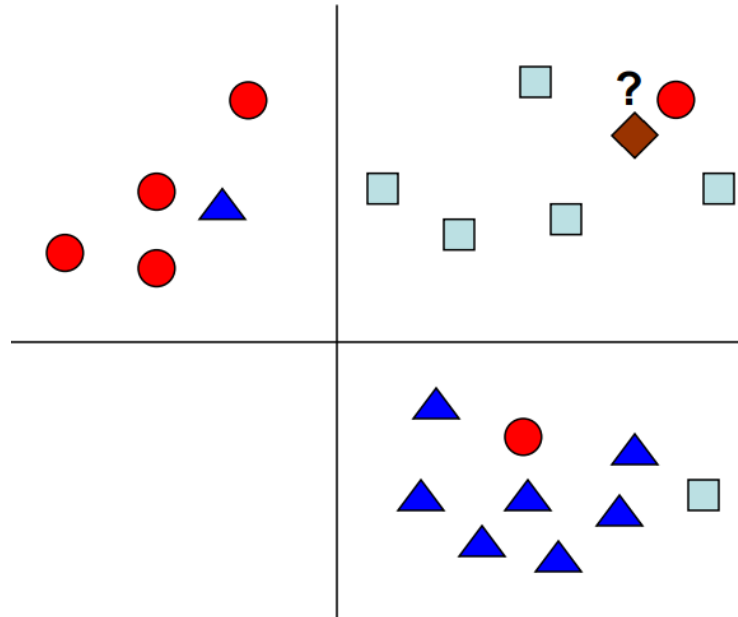
- Model observations (x, y) first ($P(x, y)$), then infer $P(y|x)$
- Good for missing variables, better diagnostics
- Easy to add prior knowledge about data

Today's Topics

- Type of classifiers
- **KNN**
- Setting Parameters
- Analysis of KNN

KNN

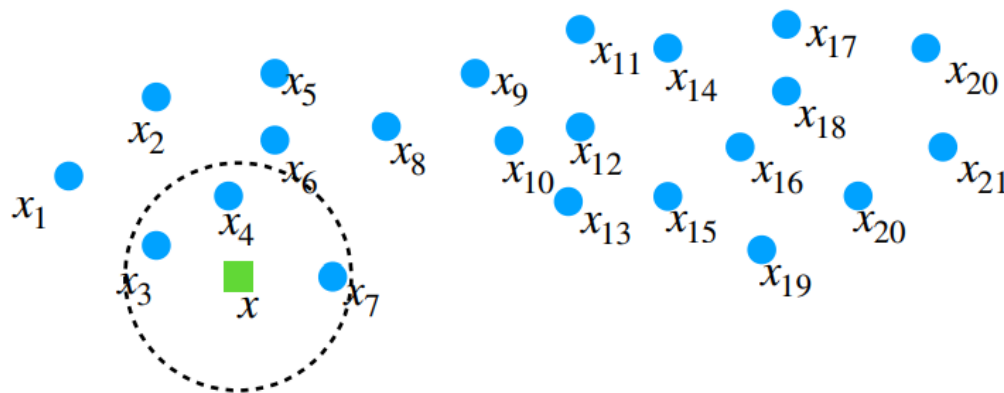
- A simple, yet surprisingly efficient algorithm
- Requires the definition of a **distance function** or similarity measures between samples
- Select the class based on the **majority vote** in the k closest points



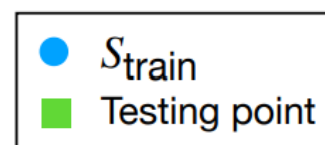
Step1: Find nearest neighbors

$$nbh_{S_{train},k}: \mathcal{X} \rightarrow \mathcal{X}^k$$

$x \mapsto \{k \text{ elements of } S_{train} \text{ which are the closest to } x\}$



How to define the distance?



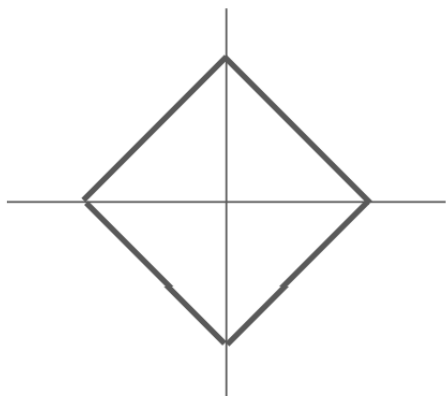
$$nbh_{S_{train},3}(x) = \{x_3, x_4, x_7\}$$

Distance Metric

- Distance Metric

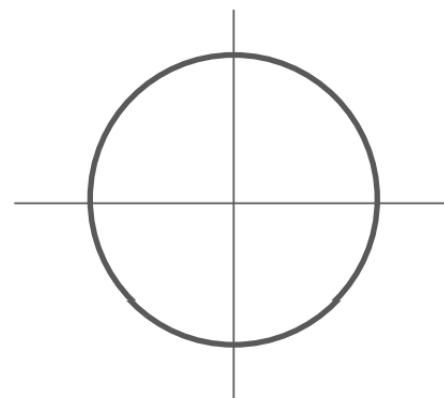
L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 (Euclidean) distance

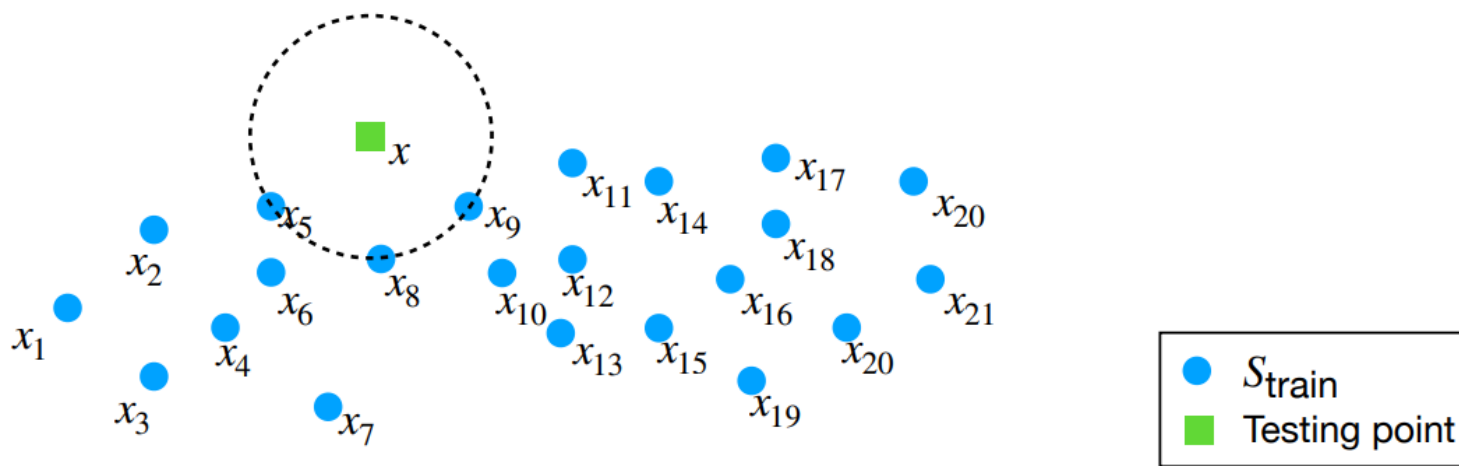
$$d_1(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



Step1: Find nearest neighbors

$$nbh_{S_{train},k}: \mathcal{X} \rightarrow \mathcal{X}^k$$

$x \mapsto \{k \text{ elements of } S_{train} \text{ which are the closest to } x\}$



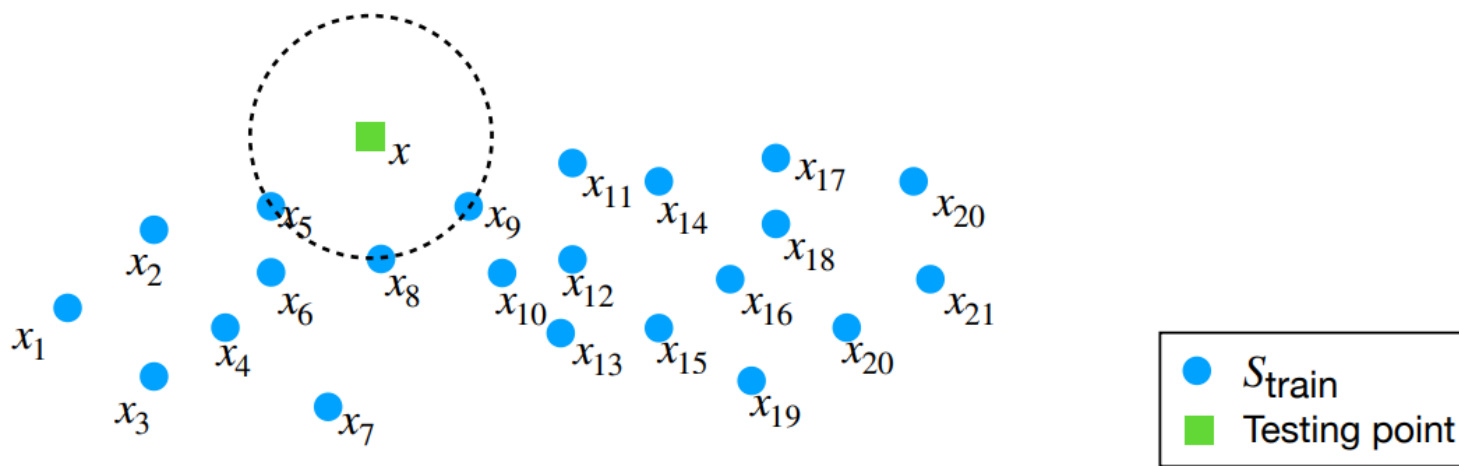
$$nbh_{S_{train},2}(x) = \{x_5, x_8\}$$

It seems that $\{x_5, x_9\}$ and $\{x_8, x_9\}$ work fine as well!

Step1: Find nearest neighbors

$$nbh_{S_{train},k}: \mathcal{X} \rightarrow \mathcal{X}^k$$

$x \mapsto \{k \text{ elements of } S_{train} \text{ which are the closest to } x\}$

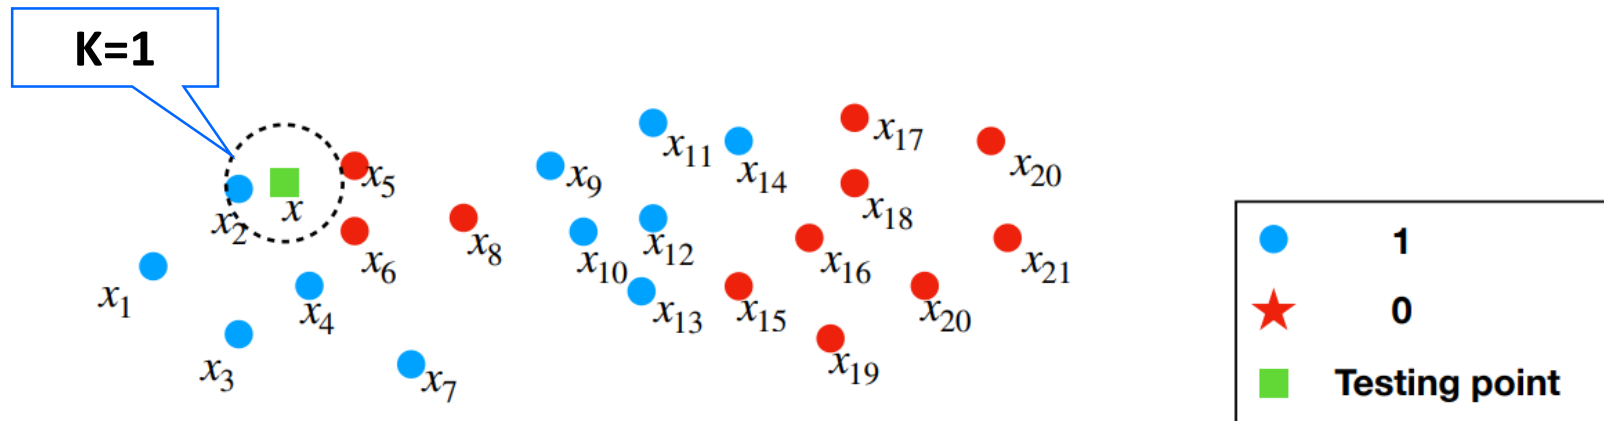


$$nbh_{S_{train},2}(x) = \{x_5, x_8\}$$

Not uniquely defined!
It will depend on the strategy
Often ties are broken **randomly**

Step2: Select Class

$$f_{S_{train},k}(x) = \text{majority}\{y_i: x_i \in nbh_{S_{train},k}(x)\}$$

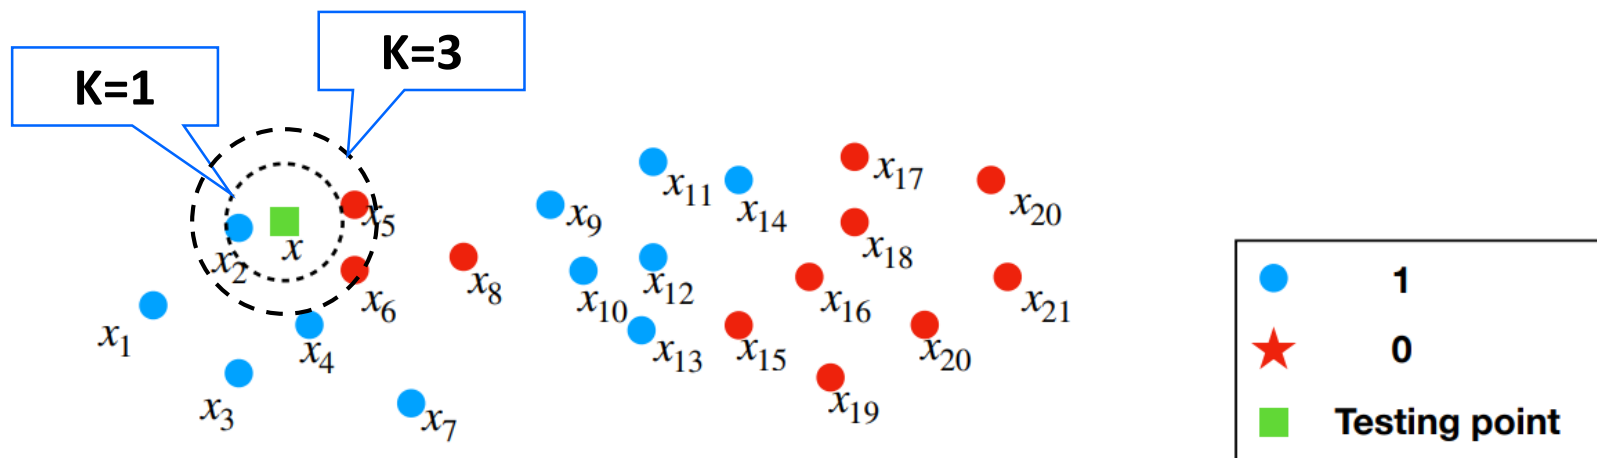


$$f_{S_{train},1}(x) = 1$$

$$f_{S_{train},3}(x) = ?$$

Step2: Select Class

$$f_{S_{train},k}(x) = \text{majority}\{y_i: x_i \in nbh_{S_{train},k}(x)\}$$

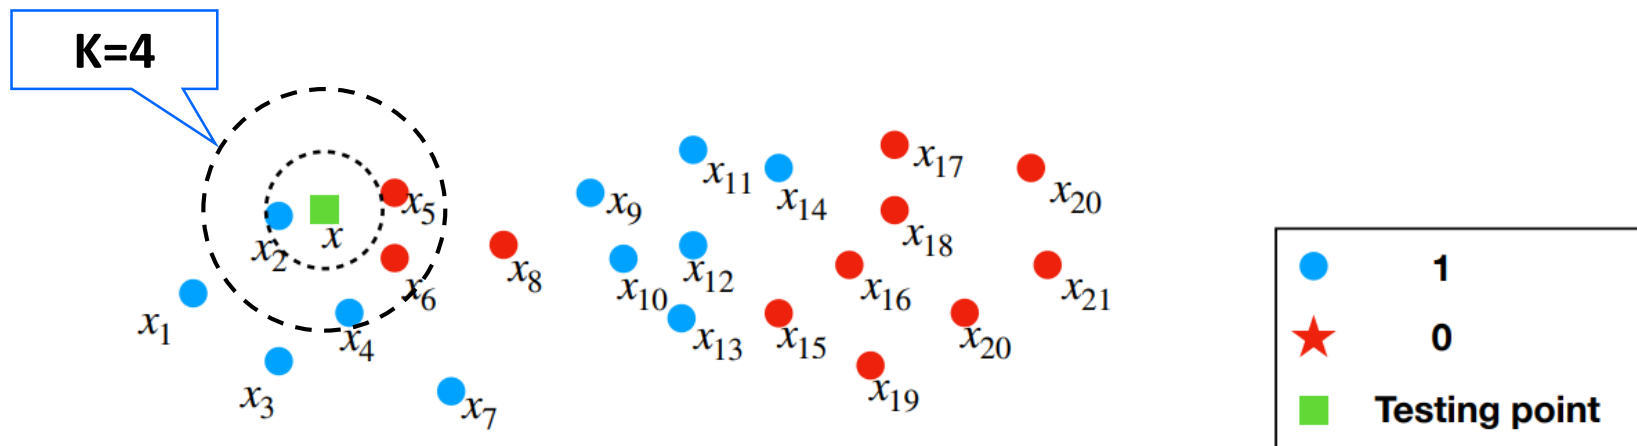


$$f_{S_{train},1}(x) = 1$$

$$f_{S_{train},3}(x) = 0$$

Step2: Select Class

$$f_{S_{train},k}(x) = \text{majority}\{y_i: x_i \in nbh_{S_{train},k}(x)\}$$



$$f_{S_{train},4}(x) = ?$$

Tie!

For the binary case it is good to pick k to be **odd** so that there is a clear winner.

KNN

- **Summary**
- **Step1: Find nearest neighbors**

L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$

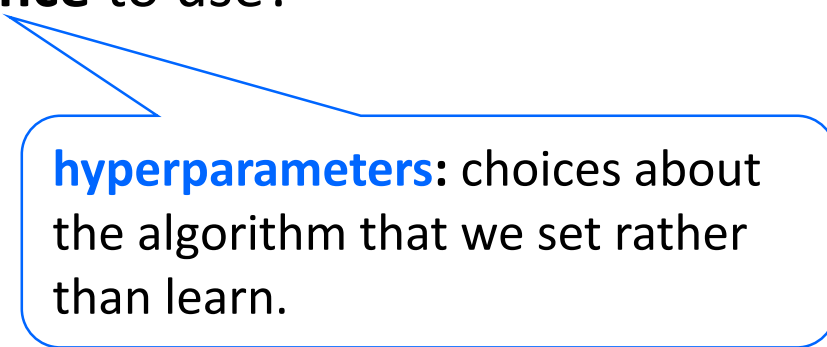
- **Step2: Select Class (majority vote)**

Today's Topics

- Type of classifiers
- KNN
- *Setting Parameters*
- Analysis of KNN

Setting Parameters

- What do we need to set for KNN?
- What is the best **value of k** to use?
- What is the best **distance** to use?



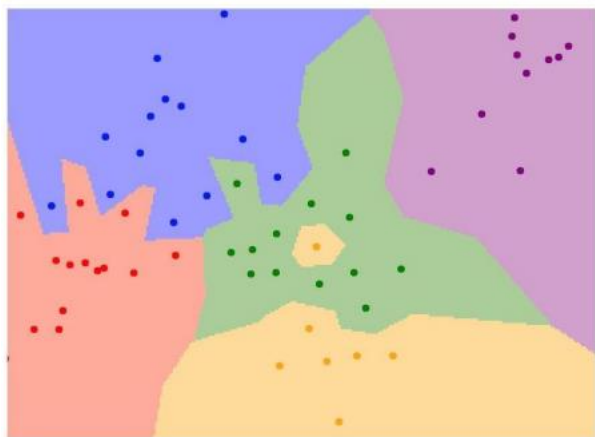
hyperparameters: choices about the algorithm that we set rather than learn.

Very problem-dependent.

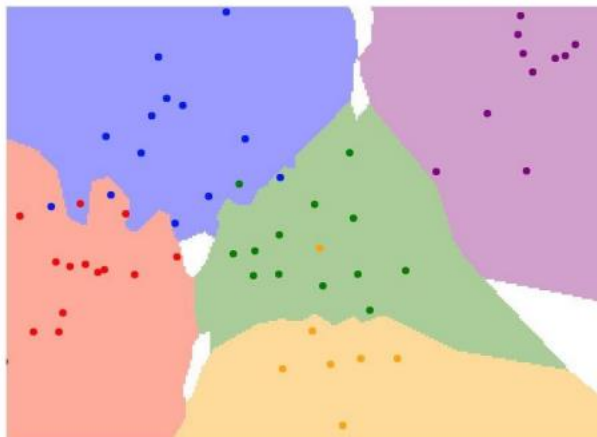
Must try them all out and see what works best.

Setting Hyperparameters

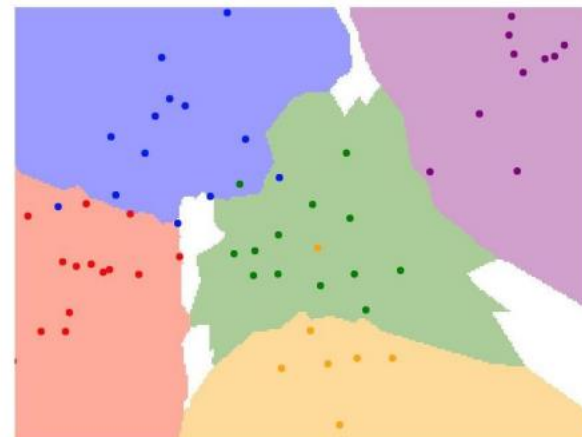
- Results in different **value of k**



$K = 1$



$K = 3$



$K = 5$

Setting Hyperparameters

- Results in different **distance metrics**

L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



K = 1

L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



K = 1

Setting Hyperparameters

Idea #1: Choose hyperparameters
that work best on the data



Your Dataset

Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data

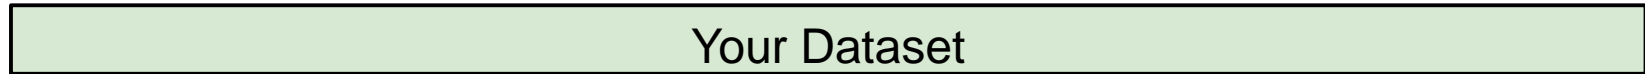


Your Dataset

Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data



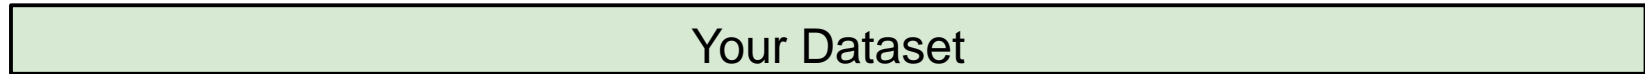
Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data



Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data



Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

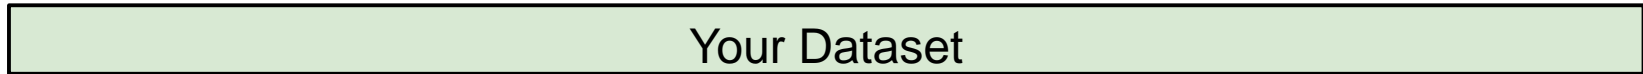
BAD: No idea how algorithm will perform on new data



Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data



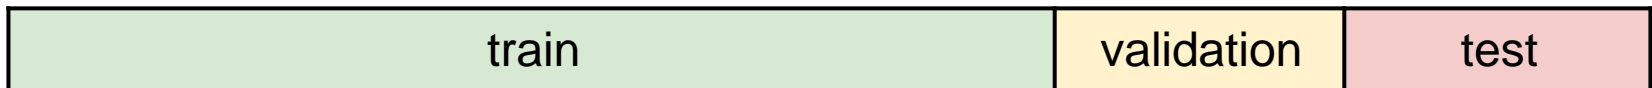
Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data



Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!



Any better solutions?

Setting Hyperparameters

Your Dataset

Idea #4: Cross-Validation: Split data into **folds**, try each fold as validation and average the results

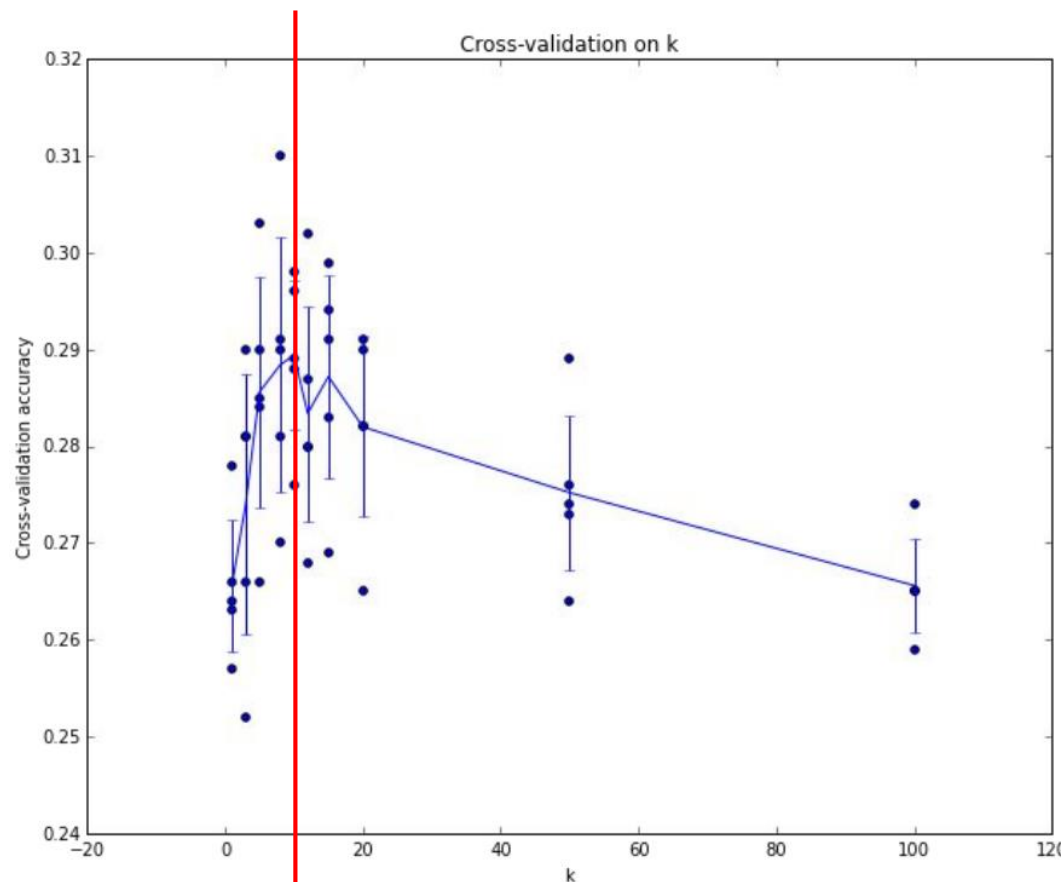
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but not used too frequently in deep learning

Why?

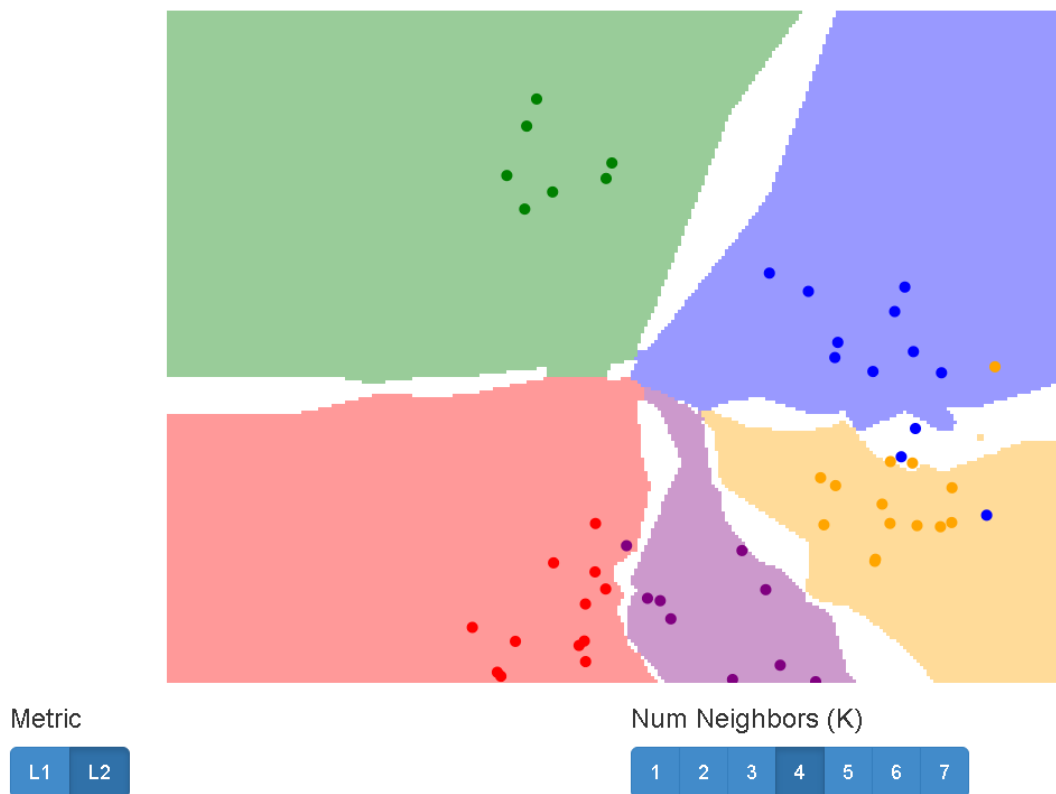
Setting Hyperparameters

- Example of **5-fold cross-validation** for the value of k .
- Each point: single outcome.
- The line goes through the mean, bars indicated standard deviation
- Seems that $k \approx 7$ works best for this data



Setting Hyperparameters

- Run the demo with different hyperparameters



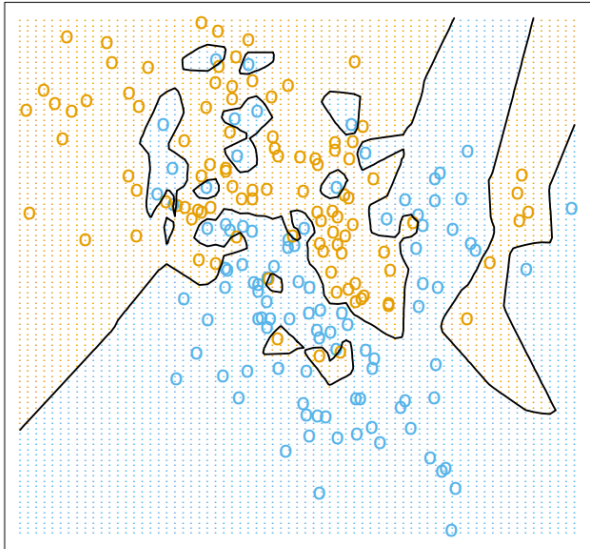
<http://vision.stanford.edu/teaching/cs231n-demos/knn/>

Today's Topics

- Type of classifiers
- KNN
- Setting Parameters
- *Analysis of KNN*

Bias-Variance for KNN

$K=1$



Small k

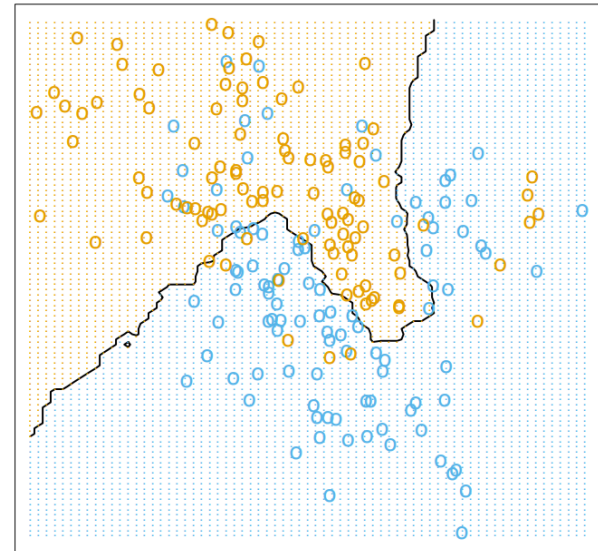
Small bias

Very complex decision boundary

Large variance

Overfitting

$K=15$



Large k

Large bias

extreme case: $k=n$, constant prediction

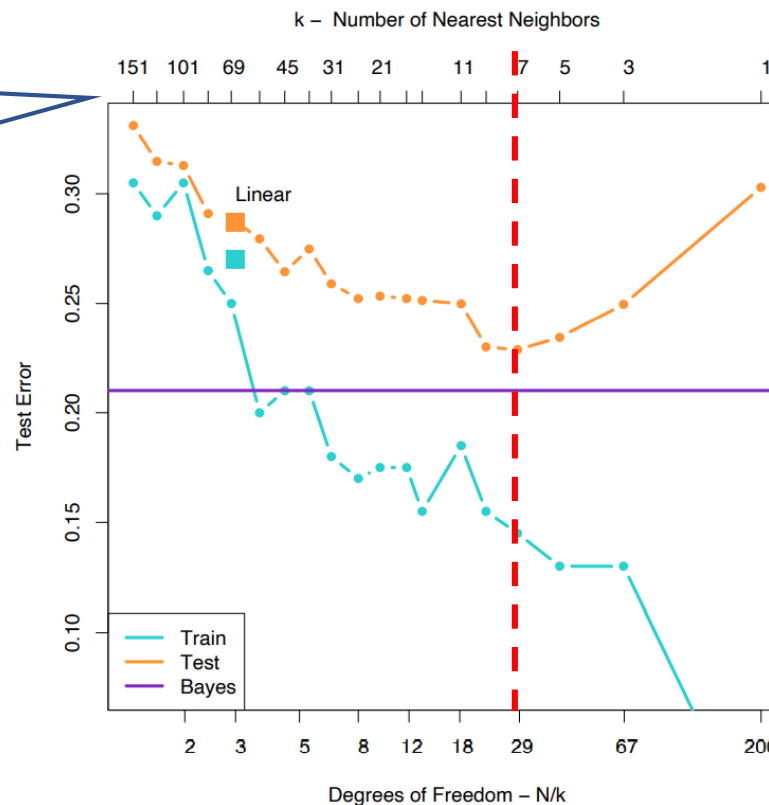
Small variance

Bias-Variance for KNN

Complexity increases when k decreases

Large k
Large bias
Small variance

Small k
Small bias
Large variance



Good k
Small bias
complex enough
decision boundary
Small variance
no overfitting

Complexity of KNN

Q: With N examples, how fast are **training** and **prediction**?

	Training	Prediction
Complexity	$O(1)$	$O(N)$



Complexity of KNN

Q: With N examples, how fast are **training** and **prediction**?

	Training	Prediction
Complexity	$O(1)$	$O(N)$
Action	Simply remembers all the training data No explicit training process <i>“Lazy Learning”</i>	For each test sample: Find closest training sample Predict label of nearest sample

Complexity of KNN

Q: With N examples, how fast are **training** and **prediction**?

	Training	Prediction
Complexity	$O(1)$	$O(N)$
Action	Simply remembers all the training data No explicit training process <i>"Lazy Learning"</i>	For each test sample: Find closest training sample Predict label of nearest sample



This is bad.



Complexity of KNN

Q: With N examples, how fast are **training** and **prediction**?

	Training	Prediction
Complexity	$O(1)$	$O(N)$
Action	Simply remembers all the training data No explicit training process <i>"Lazy Learning"</i>	For each test sample: Find closest training sample Predict label of nearest sample



This is bad.

- ◆ We want classifiers that are **fast at prediction**; slow for training is ok.
- ◆ **Test time performance** is usually much more important in practice.

Can we use KNN on images?

- Very slow at test time
- Distance metrics on pixels are not informative

Never

Original



Boxed



Shifted



Tinted

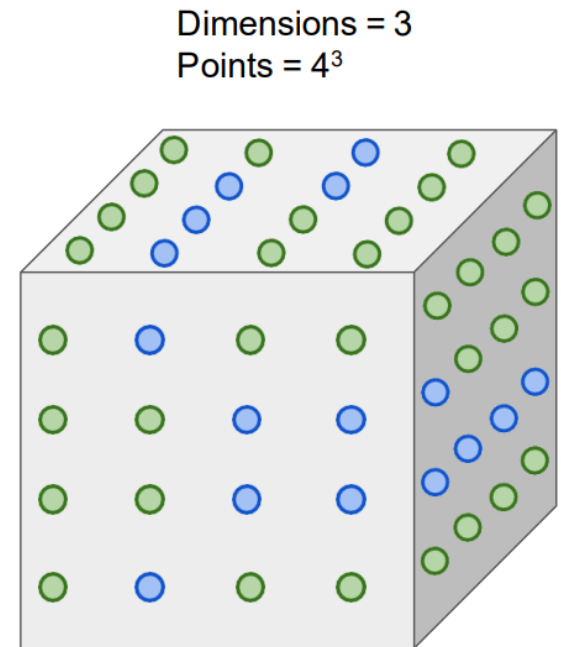
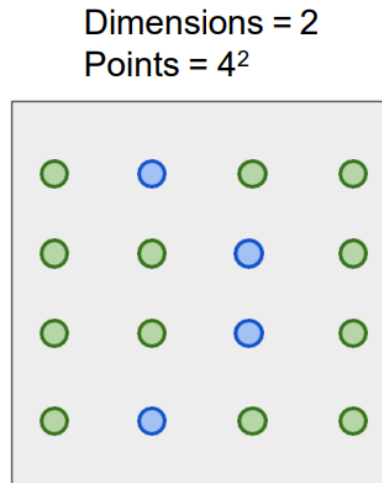
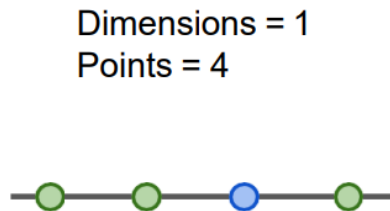


(all 3 images **have same L2 distance** to the one on the left)

Can we use KNN on images?

Never

- *Curse of dimensionality*
- In high-dimensional situations, the data samples are sparse and the distance calculation is difficult



Summary

- **KNN Algorithm**
 - Step1: Find nearest neighbors
 - Step2: Select Class (majority vote)
- **Setting Hyperparameters**
 - value of k
 - distance metric
- **Analysis of KNN**
 - bias and variance
 - complexity(train/predict)

Summary

- **Strength/Weakness of KNN**

- ✓ Simple to implement and intuitive to understand
- ✓ Can learn non-linear decision
- ✓ No Training Time
- × High prediction complexity for large datasets
- × Higher prediction complexity with higher dimension
- × KNN Assumes equal importance to all features
- × Sensitive to outliers

- **When should we use KNN?**

- spatial correlation
 - e.g. Recommender system: similarity between users can be viewed as distance)
- low dimension
 - e.g. Text mining

Practice

- When $k=1/3/5$, which class will the KNN algorithm discriminate the test sample into?

