


Machine Learning

CNN

Dr. Shuang LIANG


Recall: Neuron



Recall: Activation function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Recall: Neural Network



Recall: Backpropagation



Today's Topics

- CNN Development
- CNN Layer
- CNN Architecture

Today's Topics

- *CNN Development*
- CNN Layer
- CNN Architecture

A bit of history

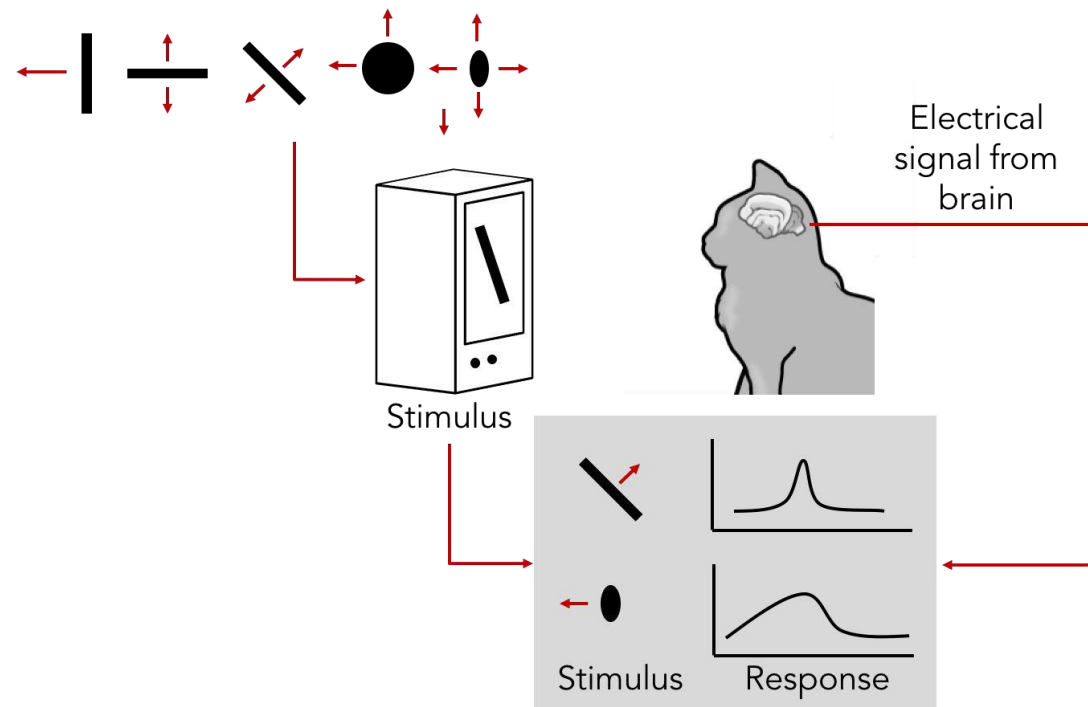
Hubel & Wiesel, 1959

Receptive fields of single
neurons in the cat's striate
cortex

1962

Receptive fields, binocular
interaction and functional
architecture in the cat's visual
cortex

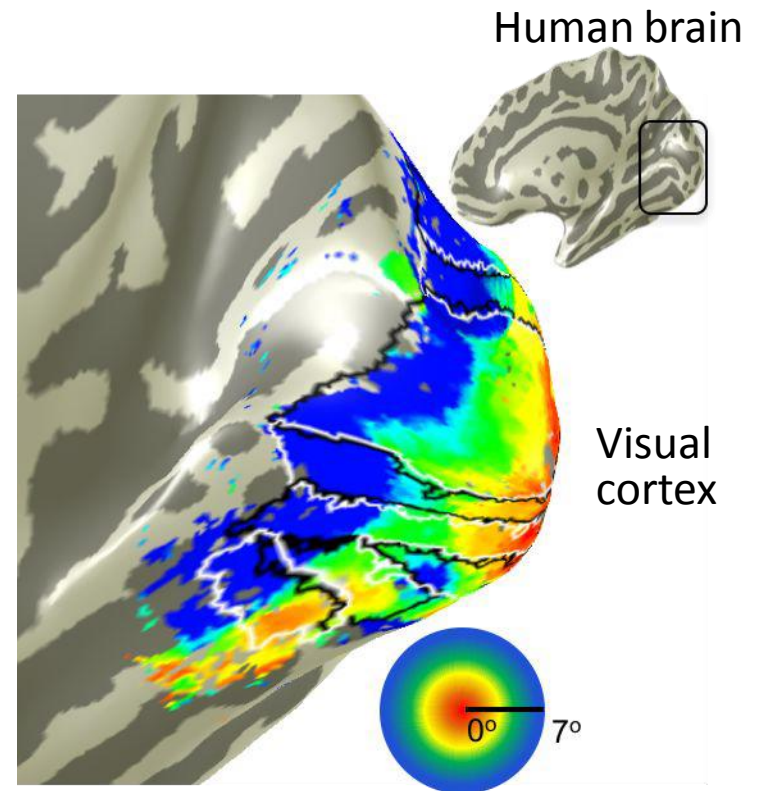

1968...



[Cat image](#) by CNX OpenStax is licensed
under CC BY 4.0; changes made

A bit of history

Topographical mapping in the cortex:
nearby cells in cortex represent
nearby regions in the visual field



Retinotopy images courtesy of Jesse Gomez in the
Stanford Vision & Perception Neuroscience Lab.

Hierarchical organization




Illustration of hierarchical organization in early visual pathways by Lane McIntosh, copyright CS231n 2017

Simple cells:
Response to light orientation

Complex cells:
Response to light orientation and movement

Hypercomplex cells:
response to movement with an end point



No response




Response
(end point)

A bit of history

Neocognitron *[Fukushima 1980]*


“sandwich” architecture (SCSCSC...)
simple cells: modifiable parameters
complex cells: perform pooling



A bit of history

Gradient-based learning applied to document recognition

[LeCun, Bottou, Bengio, Haffner 1998]



LeNet-5

A bit of history

ImageNet Classification with Deep Convolutional Neural Networks

[Krizhevsky, Sutskever, Hinton, 2012]




Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

“AlexNet”

Fast-forward to today: ConvNets are everywhere

Classification





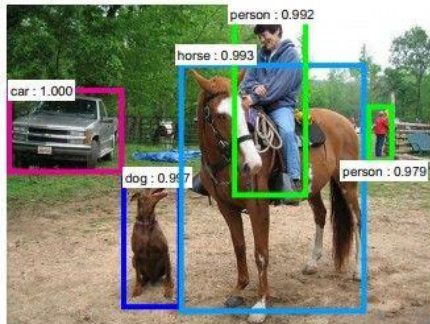
Retrieval



Figures copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012.
Reproduced with permission.

Fast-forward to today: ConvNets are everywhere


Detection



Figures copyright Shaoqing Ren, Kaiming He, Ross Girshick, Jian Sun, 2015.
Reproduced with permission.

[Faster R-CNN: Ren, He, Girshick, Sun 2015]

Segmentation



Figures copyright Clement Farabet, 2012.
Reproduced with permission.

[Farabet et al., 2012]

Fast-forward to today: ConvNets are everywhere



Photo by Lane McIntosh. Copyright CS231n 2017.

self-driving cars




[This image](#) by GBPublic_PR is licensed under [CC-BY 2.0](#)

NVIDIA Tesla line

(these are the GPUs on rye01.stanford.edu)

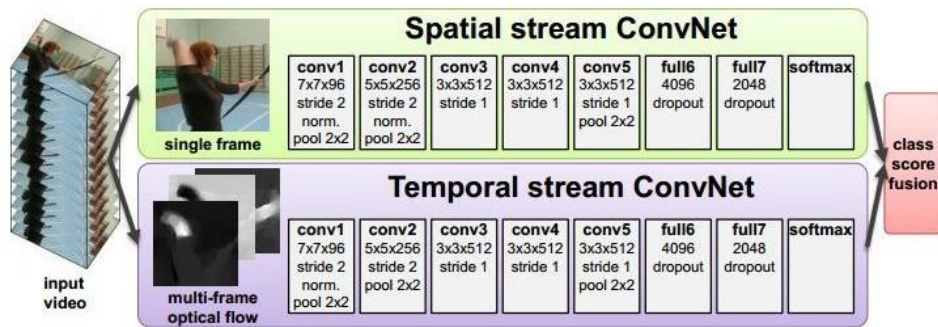
Note that for embedded systems a typical setup would involve NVIDIA Tegras, with integrated GPU and ARM-based CPU cores.

Fast-forward to today: ConvNets are everywhere



[Taigman et al. 2014]


Activations of [inception-v3 architecture](#) [Szegedy et al. 2015] to image of Emma McIntosh, used with permission. Figure and architecture not from Taigman et al. 2014.



[Simonyan et al. 2014]

Figures copyright Simonyan et al., 2014. Reproduced with permission.

Illustration by Lane McIntosh, photos of Katie Cumnock used with permission.





Fast-forward to today: ConvNets are everywhere



[Toshev, Szegedy 2014]


Images are examples of pose estimation, not actually from Toshev & Szegedy 2014. Copyright Lane McIntosh.



[Guo et al. 2014]

Figures copyright Xiaoxiao Guo, Satinder Singh, Honglak Lee, Richard Lewis, and Xiaoshi Wang, 2014. Reproduced with permission.

Fast-forward to today: ConvNets are everywhere



[Levy et al. 2016]

Figure copyright Levy et al. 2016.
Reproduced with permission.



[Dieleman et al. 2014]

From left to right:
[public domain by NASA](#),
usage [permitted](#) by ESA/Hubble,
[public domain by NASA](#),
and [public domain](#).



[Sermanet et al. 2011]
[Ciresan et al.]

Photos by Lane McIntosh.
Copyright CS231n 2017.

Fast-forward to today: ConvNets are everywhere

[This image](#) by Christin Khan is in the public domain and originally came from the U.S. NOAA.



Whale recognition, Kaggle Challenge

Photo and figure by Lane McIntosh; not actual example from Mnih and Hinton, 2010 paper.



Mnih and Hinton, 2010

Fast-forward to today: ConvNets are everywhere

No errors



A white teddy bear sitting in the grass

Minor errors



A man in a baseball uniform throwing a ball

Somewhat related



A woman is holding a cat in her hand

Image Captioning

[Vinyals et al., 2015]

[Karpathy and Fei-Fei, 2015]



A man riding a wave on top of a surfboard



A cat sitting on a suitcase on the floor



A woman standing on a beach holding a surfboard

All images are CC0 Public domain:

<https://pixabay.com/en/luggage-antique-cat-1643010/>

<https://pixabay.com/en/teddy-plush-bears-cute-teddy-bear-1623436/>

<https://pixabay.com/en/surf-wave-summer-sport-litoral-1668716/>



<https://pixabay.com/en/woman-female-model-portrait-adult-983967/>

<https://pixabay.com/en/handstand-lake-meditation-496008/>


<https://pixabay.com/en/baseball-player-shortstop-infield-1045263/>

Captions generated by Justin Johnson using [NeuralTalk2](#)

Fast-forward to today: ConvNets are everywhere



Figures copyright Justin Johnson, 2015. Reproduced with permission. Generated using the Inceptionism approach from a [blog post](#) by Google Research.



[Original image](#) is CC0 public domain

[Starry Night](#) and [Tree Roots](#) by Van Gogh are in the public domain
[Bokeh image](#) is in the public domain

Stylized images copyright Justin Johnson, 2017; reproduced with permission



Gatys et al, "Image Style Transfer using Convolutional Neural Networks", CVPR 2016

Gatys et al, "Controlling Perceptual Factors in Neural Style Transfer", CVPR 2017

Today's Topics


- CNN Development
- *CNN Layer*
- CNN Architecture

CNN Layer

- *Fully Connected Layer*
- Convolution Layer
- Pooling Layer

Fully Connected Layer

32x32x3 image -> stretch to 3072 x 1




CNN Layer

- Fully Connected Layer
- *Convolution Layer*
- Pooling Layer


Convolution Layer

32x32x3 image -> preserve spatial structure



Convolution Layer

32x32x3 image

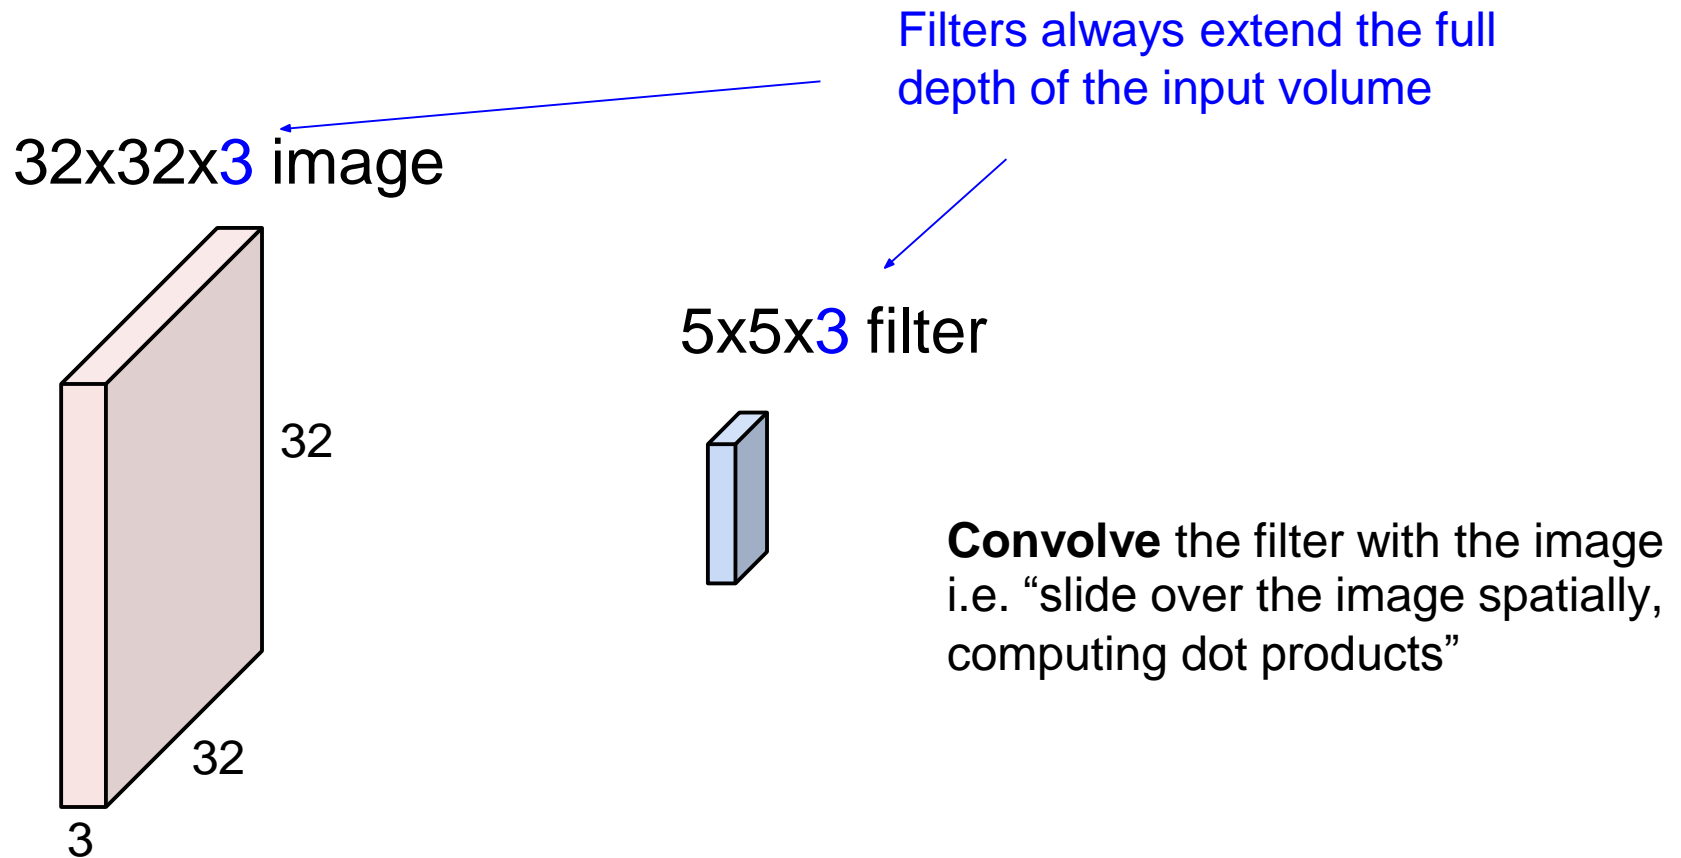


5x5x3 filter

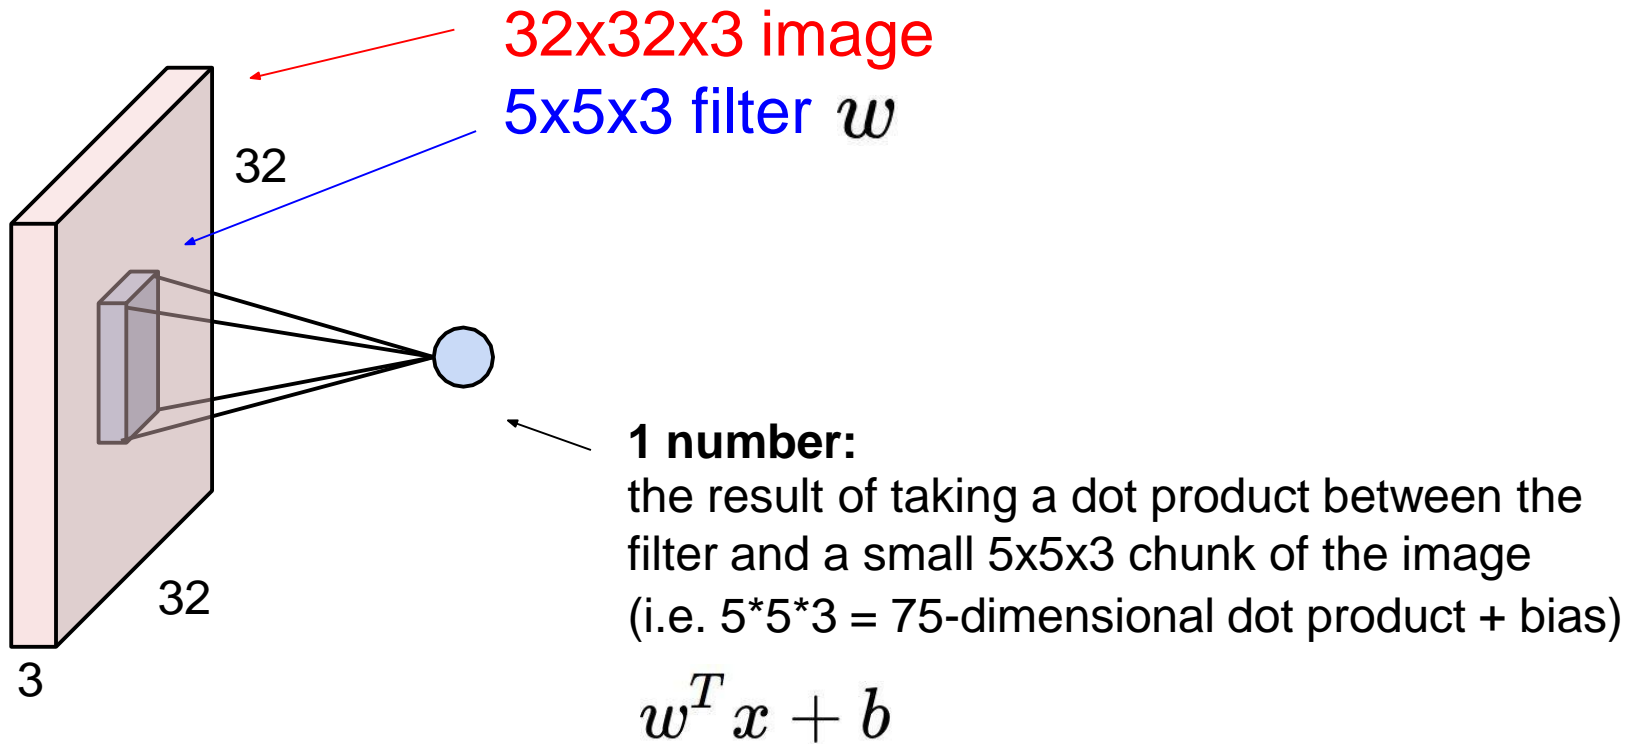


Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”


Convolution Layer



Convolution Layer




Convolution Layer




Convolution Layer

consider a second, **green** filter



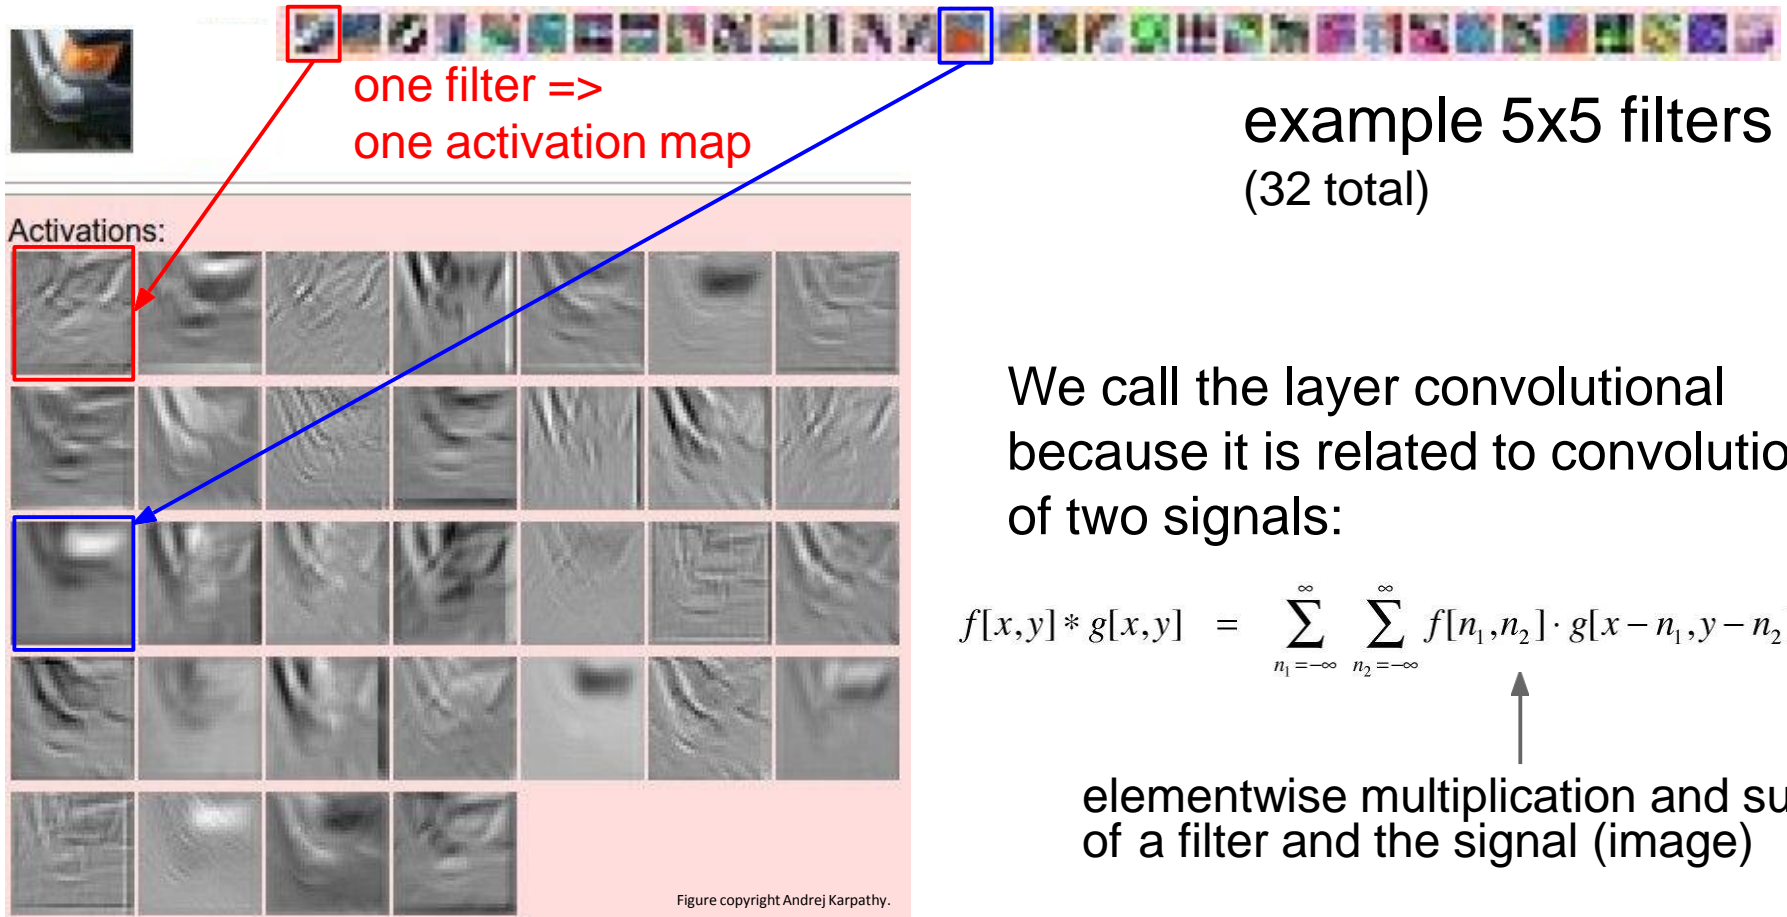
Convolution Layer

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!

Convolution Layer



Convolution


stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



Convolution


If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



We set stride=1 below

Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1


Filter 1

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

A closer look at spatial dimensions




A closer look at spatial dimensions




7x7 input (spatially)
assume 3x3 filter

A closer look at spatial dimensions




7x7 input (spatially)
assume 3x3 filter

A closer look at spatial dimensions




7x7 input (spatially)
assume 3x3 filter

A closer look at spatial dimensions



7x7 input (spatially)
assume 3x3 filter


A closer look at spatial dimensions



7x7 input (spatially)
assume 3x3 filter

7 **=> 5x5 output**

A closer look at spatial dimensions




7x7 input (spatially)

assume 3x3 filter

Applied **with stride 2**

A closer look at spatial dimensions




7x7 input (spatially)

assume 3x3 filter

Applied **with stride 2**

A closer look at spatial dimensions




7

7x7 input (spatially)
assume 3x3 filter

Applied **with stride 2**
=> 3x3 output!

A closer look at spatial dimensions



7x7 input (spatially)


assume 3x3 filter

Applied **with stride 3?**

doesn't fit!

cannot apply 3x3 filter on
7x7 input with stride 3.

Convolution Layer



Output size:
 $(N - F) / \text{stride} + 1$

e.g. $N = 7$, $F = 3$:

stride 1 $\Rightarrow (7 - 3) / 1 + 1 = 5$

stride 2 $\Rightarrow (7 - 3) / 2 + 1 = 3$

stride 3 $\Rightarrow (7 - 3) / 3 + 1 = 2.33 \therefore \backslash$

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

(recall:)

$$(N - F) / \text{stride} + 1$$

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!

in general, common to see CONV layers with stride 1, filters of size $F \times F$, and zero-padding with $(F-1)/2$. (will preserve size spatially)


e.g. $F = 3 \Rightarrow$ zero pad with 1

$F = 5 \Rightarrow$ zero pad with 2

$F = 7 \Rightarrow$ zero pad with 3

Convolution Layer: Example


32x32 input convolved repeatedly with 5x5 filters shrinks volumes spatially! (32 \rightarrow 28 \rightarrow 24 ...). Shrinking too fast is not good, doesn't work well.



Convolution Layer: Example

Input volume: **32x32x3**
10 5x5 filters with stride 1, pad 2


Output volume size:?



Convolution Layer: Example

Input volume: **32x32x3**
10 **5x5** filters with stride **1**, pad **2**

Output volume size:
 $(32 + 2 * 2 - 5) / 1 + 1 = 32$ spatially, so
32x32x10




Convolution Layer: Example

Input volume: **32x32x3**

10 5x5 filters with stride 1, pad 2


Number of parameters in this layer?



Convolution Layer: Example

Input volume: **32x32x3**
10 **5x5** filters with stride 1, pad 2

Number of parameters in this layer?
each filter has $5*5*3 + 1 = 76$ params
 $\Rightarrow 76*10 = 760$ (+1 for bias)



Convolution Layer: Summary

- Accepts a volume of size $W_1 \times H_1 \times D_1$

- Requires 4 hyperparameters:

- Numbers of filters K
- Their spatial extent F
- The stride S
- The amount of zero padding P

Common settings:

K = (powers of 2, e.g. 32, 64, 128, 512)

- $F = 3, S = 1, P = 1$
- $F = 5, S = 1, P = 2$
- $F = 5, S = 2, P = ?$ (whatever fits)
- $F = 1, S = 1, P = 0$

- Produces a volume of size $W_2 \times H_2 \times D_2$ where:

- $W_2 = \frac{W_1 - F + 2P}{S} + 1$

- $H_2 = \frac{H_1 - F + 2P}{S} + 1$ (i.e. width and height are computed equally by symmetry)

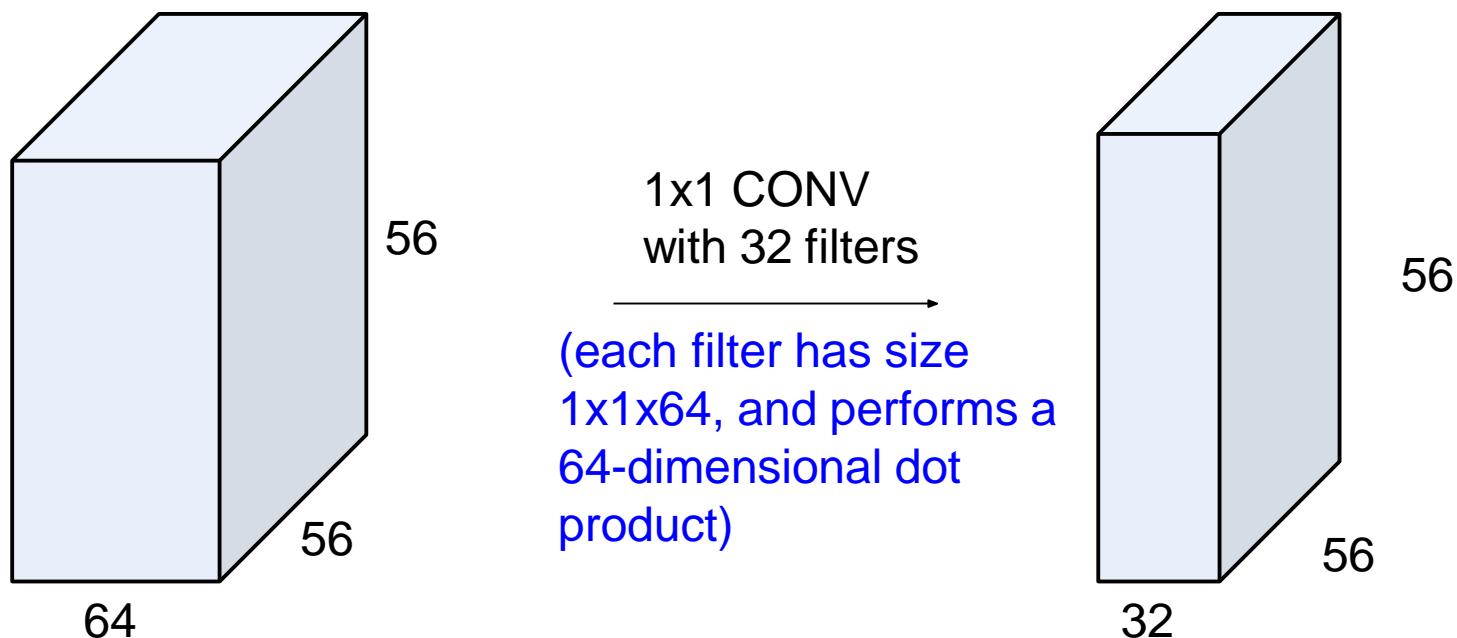
- $D_2 = K$

Convolution Layer: Summary

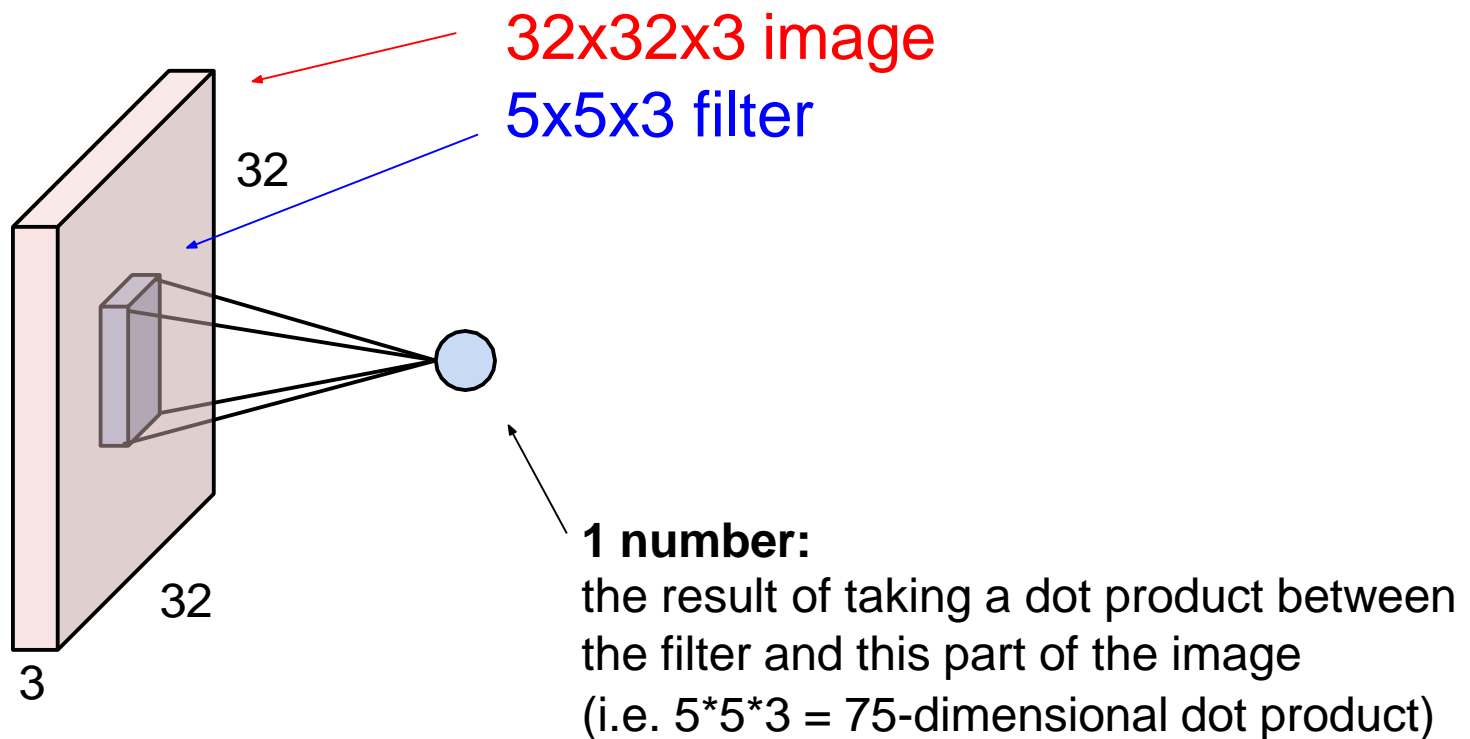
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and the offset by d -th bias.

1×1 conv


- 1x1 convolution layers make perfect sense



The brain/neuron view of Conv Layer




The brain/neuron view of Conv Layer





1 number:

the result of taking a dot product between the filter and this part of the image (i.e. $5 \times 5 \times 3 = 75$ -dimensional dot product)



It's just a neuron with local connectivity...

The brain/neuron view of Conv Layer





An activation map is a 28x28 sheet of neuron outputs:

1. Each is connected to a small region in the input
2. All of them share parameters

“5x5 filter” -> “5x5 receptive field for each neuron”

The brain/neuron view of Conv Layer




E.g. with 5 filters,
CONV layer consists of
neurons arranged in a 3D grid
(28x28x5)

There will be 5 different
neurons all looking at the same
region in the input volume

Reminder: FC Layer

32x32x3 image -> stretch to 3072 x 1

Each neuron
looks at the full
input volume




CNN Layer


- Fully Connected Layer
- Convolution Layer
- *Pooling Layer*

Pooling Layer

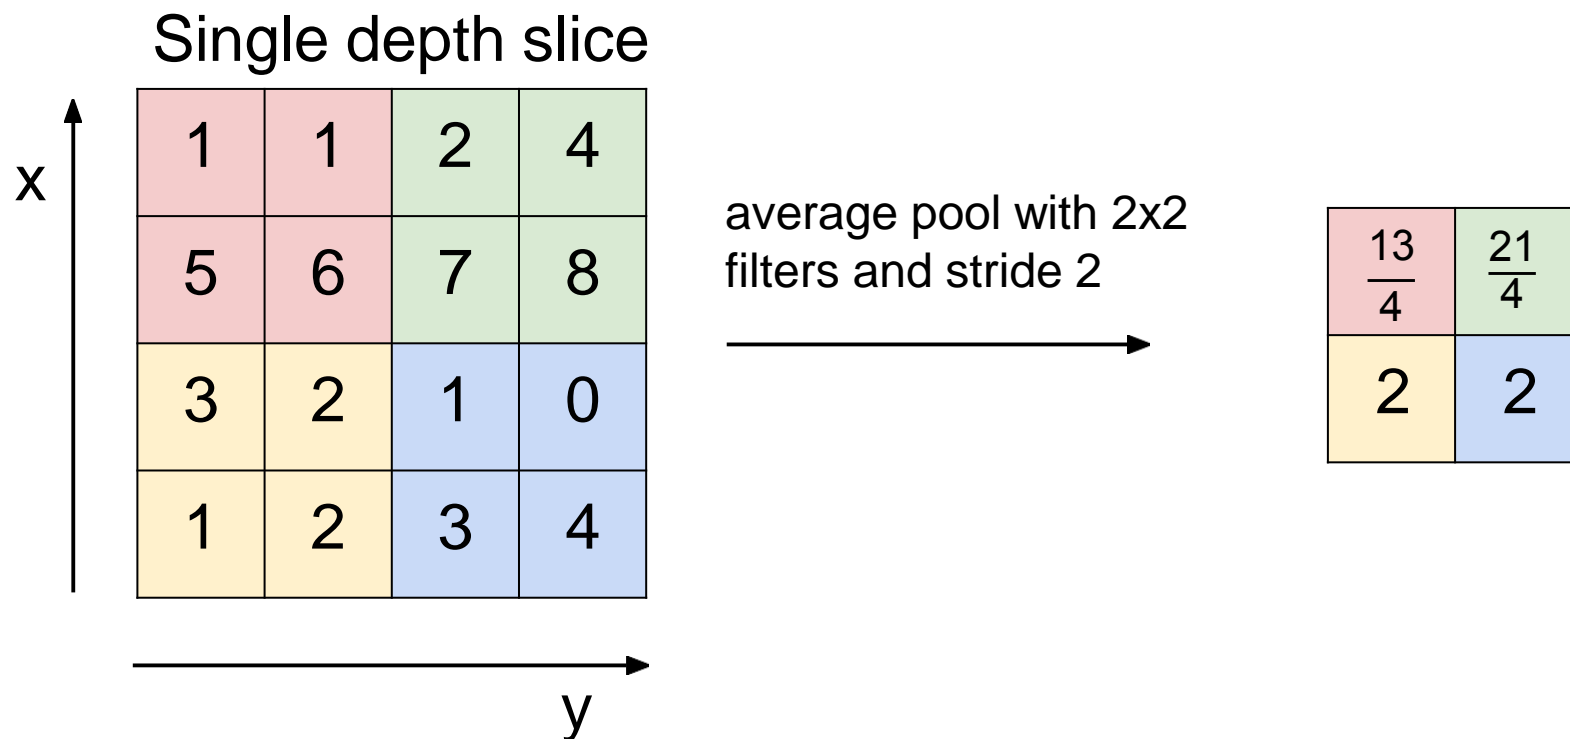
- makes the representations smaller and more manageable
- operates over each activation map independently:



Max Pooling



Average Pooling



Pooling Layer: Summary

Common settings:

$F = 2, S = 2$

$F = 3, S = 2$


- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires 2 hyperparameters:
 - Their spatial extent F
 - The stride S
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = \frac{W_1 - F}{S} + 1$
 - $H_2 = \frac{H_1 - F}{S} + 1$
 - $D_2 = D_1$
- Introduces zero parameters since it computes a fixed function of the input
- Note that it is not common to use zero-padding for Pooling layers

Today's Topics

- CNN Development
- CNN Layer
- ***CNN Architecture***


CNN Architecture

- ConvNet is a sequence of Convolution Layers, interspersed with activation functions



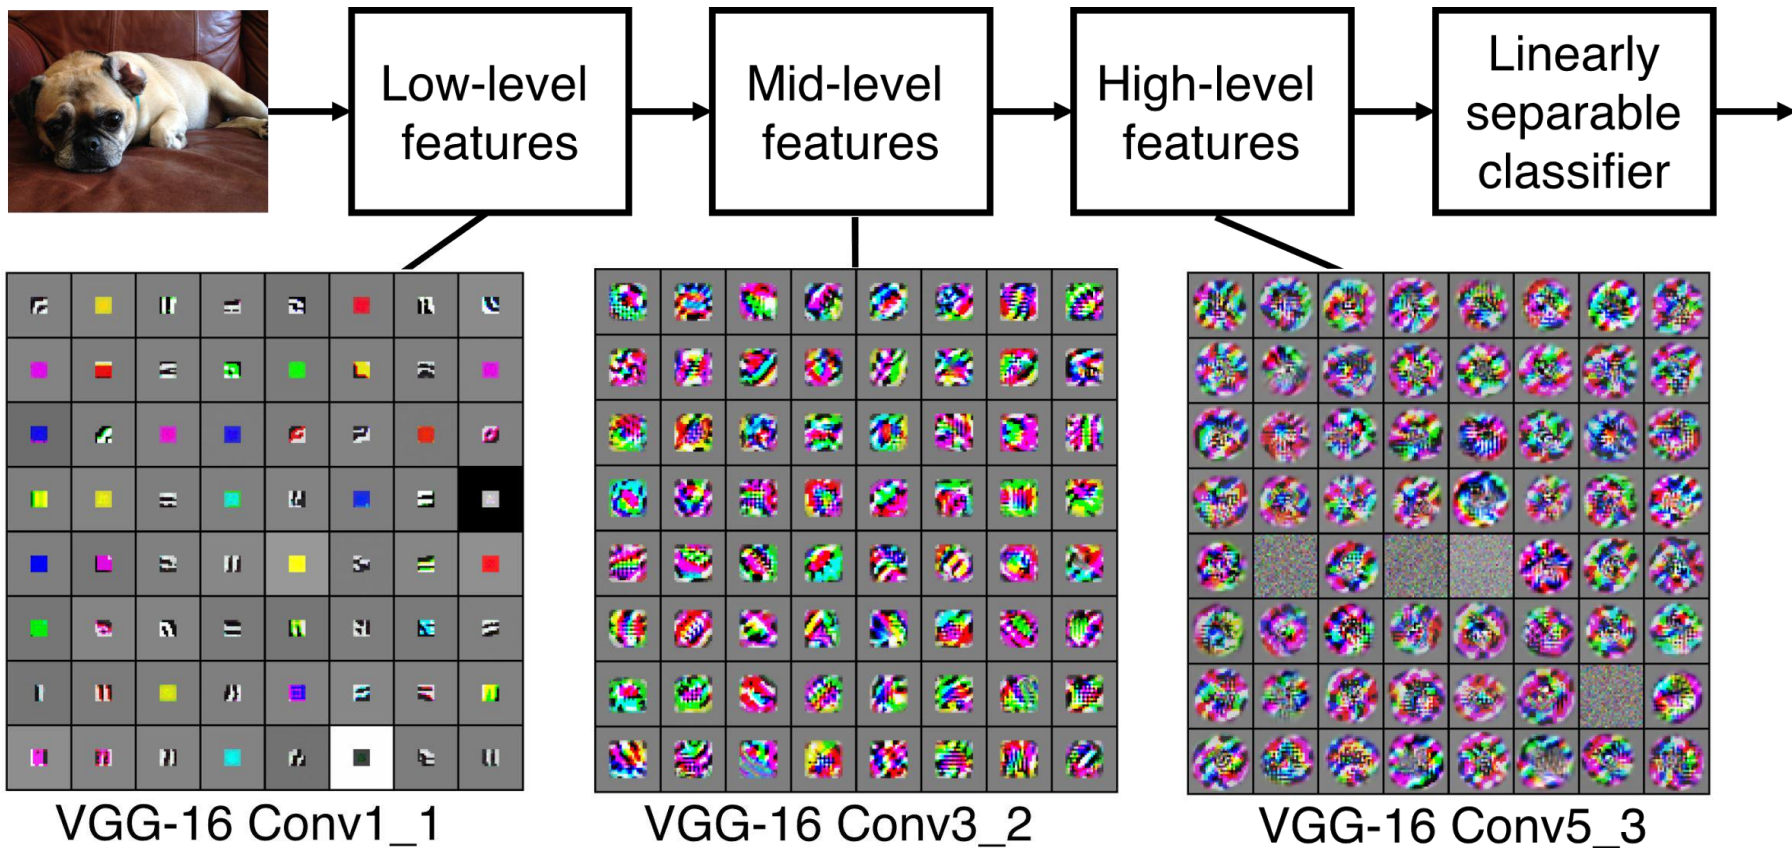
CNN Architecture

- ConvNet is a sequence of Convolution Layers, interspersed with activation functions



CNN Architecture

- ConvNet is a sequence of Convolution Layers, interspersed with activation functions



CNN Architecture

- ConvNet is a sequence of Convolution Layers, interspersed with activation functions



CNN Summary

- ConvNets stack CONV, POOL, FC layers
- Trend towards smaller filters and deeper architectures
- Trend towards getting rid of POOL/FC layers (just CONV)
- Typical architectures look like
 - **$[(\text{CONV-RELU}) * N - \text{POOL?}] * M - (\text{FC-RELU}) * K, \text{SOFTMAX}$**
 - where N is usually up to ~ 5 , M is large, $0 \leq K \leq 2$.
 - but recent advances such as ResNet/GoogLeNet challenge this paradigm

Thinking

- Why do we introduce a max pooling layer instead of using a convolutional layer with a larger stride?
- What are the advantages of convolutional layers over fully connected layers for image classification?